

DATA1030_FinalProject

December 6, 2021

```
[19]: import numpy as np
import pandas as pd
import matplotlib as mpl
# mpl.rcParams['figure.dpi'] = 250
import matplotlib
from matplotlib import pylab as plt
import joblib
```

```
[20]: import os

cwd = os.getcwd()
direct = os.path.abspath(os.path.join(cwd, os.pardir))
```

```
[21]: df = pd.read_csv(direct + '/data/WA_Fn-UseC_-Telco-Customer-Churn.csv')

# Exclude Customer ID
df = df.loc[:, df.columns != 'customerID']
```

```
[22]: # number of rows
print(df.shape[0])

# number of columns
print(df.shape[1])
```

7043

20

```
[23]: ## Change the type of Feature TotalCharges
df.TotalCharges = pd.to_numeric(df.TotalCharges, errors='coerce')
print(df['TotalCharges'].describe())
```

count	7032.000000
mean	2283.300441
std	2266.771362
min	18.800000
25%	401.450000
50%	1397.475000
75%	3794.737500

```
max      8684.800000
Name: TotalCharges, dtype: float64
```

1 1. Exploratory Data Analysis

```
[24]: # Target Variable

print(df['Churn'].value_counts())

fig = plt.figure()
colors = ['lightseagreen', 'lightcoral']
plt.rcParams["font.family"] = 'serif'
ax = (df['Churn'].value_counts()*100.0 / len(df)).plot(kind='bar', stacked =   
→ True, rot = 0, color = colors, figsize = (8,6))

ax.set_ylabel('Customers Percentage', size = 14)
ax.set_xlabel('Churn', size = 14)
ax.set_title('Churn Rate', size = 14)

# create a list to collect the plt.patches data
totals = []

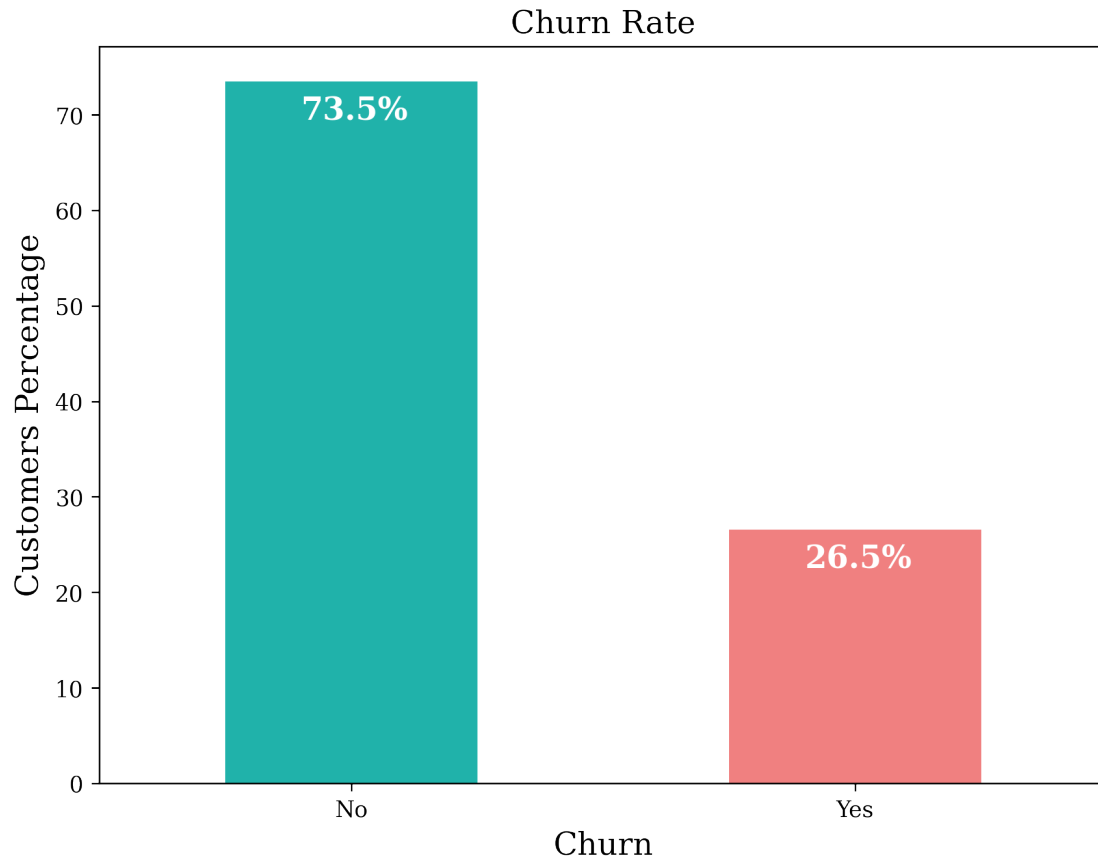
# find the values and append to list
for i in ax.patches:
    totals.append(i.get_width())

# set individual bar labels using above list
total = sum(totals)

for i in ax.patches:
    # get_width pulls left or right; get_y pushes up or down
    ax.text(i.get_x()+.15, i.get_height()-4.0, \
            str(round((i.get_height()/total), 1))+'%',
            fontsize=12,
            color='white',
            weight = 'bold',
            size = 14)


plt.savefig(direct + '/figures/Churn_Rate.png', bbox_inches='tight')
```

```
No      5174
Yes      1869
Name: Churn, dtype: int64
```



[921]: *# Churn by Contract Type*


```
count_matrix = df.groupby(['Contract', 'Churn']).size().unstack()
count_matrix_norm = count_matrix.div(count_matrix.sum(axis=1), axis=0)
count_matrix_norm.plot(kind='bar', stacked=True)
plt.ylabel('Fraction of People in Group')
plt.xlabel('Contract Type')
plt.legend(loc=4)
plt.savefig(direct + '/figures/ContractType_Churn.png',
            bbox_inches='tight', dpi=300)
plt.show()
```



```
[924]: df['tenure'].describe()

categories = df['Churn'].unique()
bin_range = (df['tenure'].min(), df['tenure'].max())


for c in categories:
    plt.hist(df[df['Churn']==c]['tenure'], alpha=0.
    ↪5, label=c, range=bin_range, bins=35, density=True)
plt.legend()
plt.ylabel('Tenure Fraction')
plt.xlabel('Churn')
plt.savefig('figures/tenure_Churn.png', bbox_inches='tight', dpi=300)
plt.show()
```



```
[925]: dataset = [df[df['Churn']=='Yes']['tenure'].values,
                  df[df['Churn']=='No']['tenure'].values]

plt.violinplot(dataset = dataset)
plt.xticks([1,2],['Left','Stayed'])
plt.ylabel('Number of Months')
plt.title("Violin Plot of Customer's Time with Platform by Customer Churn")
plt.savefig(direct + '/figures/Violin_tenure_Churn.png',
            bbox_inches='tight',dpi=300)
plt.show()
```

Violin Plot of Customer's Time with Platform by Customer Churn



```
[929]: colors = ['#4D3425', '#E4512B']


df['MonthlyCharges'].describe()


df[['MonthlyCharges', 'Churn']].boxplot(by='Churn', widths=(0.5, 0.5),
                                         boxprops=dict(color=colors[1]),
                                         <-- capprops=dict(color=colors[1]), whiskerprops=dict(color=colors[1]))
plt.ylabel('Monthly Charges ($)')
plt.xlabel('Churn')
plt.xticks([1, 2], ['Stayed', 'Left'])
plt.suptitle('')
plt.title('Monthly Charges Comparison Based on Churn')
plt.savefig('figures/boxplot_MonthlyCharges_Churn.png',
           <-- bbox_inches='tight', dpi=300)
plt.show()

categories = df['Churn'].unique()
bin_range = (df['MonthlyCharges'].min(), df['tenure'].max())

for c in categories:
    plt.hist(df[df['Churn']==c]['MonthlyCharges'], alpha=0.
    <-- 5, label=c, range=bin_range, bins=35, density=True)
plt.legend()
plt.ylabel('Monthly Charges ($)')
plt.xlabel('Churn')
plt.suptitle('')
```


```
plt.title('Monthly Charges Comparison Based on Churn')
plt.savefig(direct + '/figures/hstgm_MonthlyCharges_Churn.png',
            bbox_inches='tight', dpi=300)
plt.show()
```





[930]: *# Churn by Payment Method*

```
count_matrix = df.groupby(['PaymentMethod', 'Churn']).size().unstack()
count_matrix_norm = count_matrix.div(count_matrix.sum(axis=1), axis=0)
count_matrix_norm.plot(kind='bar', stacked=True)
plt.ylabel('Fraction of People in Each Group')
plt.xlabel('Payment Method')
plt.title('Payment Method Group Fraction Based on Churn')
plt.legend(loc=4)
plt.savefig('figures/paymentMethod_Churn.png',
            bbox_inches='tight', dpi=300)
plt.show()
```

```
[18]: # If the feature is categorical, make bar graph, and print out values counts in_
      ↪percentage and numbers
      # If the feature is continuous/ numerical, make box plot, and print out_
      ↪description

for col in df.columns:

    # categorical vs. continuous
    if df[col].dtypes == "float64":
        df[['Churn', col]].boxplot(by='Churn')
        plt.ylabel(col)
```

```


plt.suptitle('')
plt.title(col + ' Grouped by Churn Category')
plt.show()
print(df[col].describe())
df[col].plot.hist(bins = int(np.sqrt(df.shape[0])))
plt.xlabel(col)
plt.ylabel('Count')
plt.show()

# categorical vs. categorical
elif (df[col].dtypes == "object") & (col != 'Churn'):
    count_matrix = df.groupby(['Churn', col]).size().unstack()
    count_matrix_norm = count_matrix.div(count_matrix.sum(axis=1),axis=0)
    count_matrix_norm.plot(kind='bar', stacked=True)
    plt.ylabel('fraction of people in group')
    plt.legend(loc=4)
    plt.show()
    print(df[col].value_counts())
    print(df[col].value_counts(normalize=True))


# categorical vs. cateogrical
elif (df[col].dtypes == "int64") & (len(df[col].value_counts()) <= 15):
    count_matrix = df.groupby(['Churn', col]).size().unstack()
    count_matrix_norm = count_matrix.div(count_matrix.sum(axis=1),axis=0)
    count_matrix_norm.plot(kind='bar', stacked=True)
    plt.ylabel('fraction of people in group')
    plt.legend(loc=4)
    plt.show()
    print(df[col].value_counts())
    print(df[col].value_counts(normalize=True))

# cateogrical vs. continuous
elif (df[col].dtypes == "int64") & (len(df[col].value_counts()) > 15):
    df[['Churn', col]].boxplot(by='Churn')
    plt.ylabel(col)
    plt.suptitle('')
    plt.title(col + ' Grouped by Churn Category')
    plt.show()
    print(df[col].describe())
    df[col].plot.hist(bins = int(np.sqrt(df.shape[0])))
    plt.xlabel(col)
    plt.ylabel('Count')
    plt.show()


```




```
Male      3555
Female    3488
Name: gender, dtype: int64
Male      0.504756
Female    0.495244
Name: gender, dtype: float64
```




```
0    5901
1    1142
Name: SeniorCitizen, dtype: int64
0    0.837853
1    0.162147
Name: SeniorCitizen, dtype: float64
```





```
No      3641
Yes      3402
Name: Partner, dtype: int64
No      0.516967
Yes      0.483033
Name: Partner, dtype: float64
```



```
No      4933
Yes      2110
Name: Dependents, dtype: int64
No      0.700412
Yes      0.299588
Name: Dependents, dtype: float64
```



```
count    7043.000000
mean      32.371149
std       24.559481
min        0.000000
25%        9.000000
50%       29.000000
75%       55.000000
max       72.000000
Name: tenure, dtype: float64
```

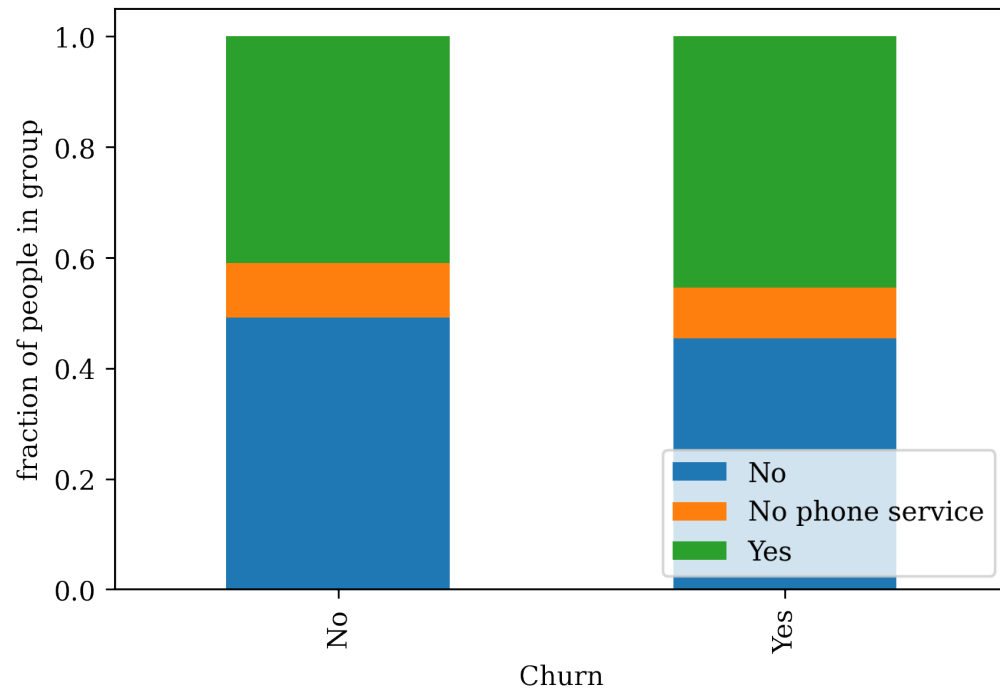


Yes 6361
No 682


```

Name: PhoneService, dtype: int64
Yes    0.903166
No     0.096834
Name: PhoneService, dtype: float64


```



```

No          3390
Yes         2971
No phone service  682
Name: MultipleLines, dtype: int64
No          0.481329
Yes         0.421837
No phone service  0.096834
Name: MultipleLines, dtype: float64


```



```

Fiber optic    3096
DSL            2421
No             1526
Name: InternetService, dtype: int64
Fiber optic    0.439585
DSL            0.343746
No             0.216669
Name: InternetService, dtype: float64

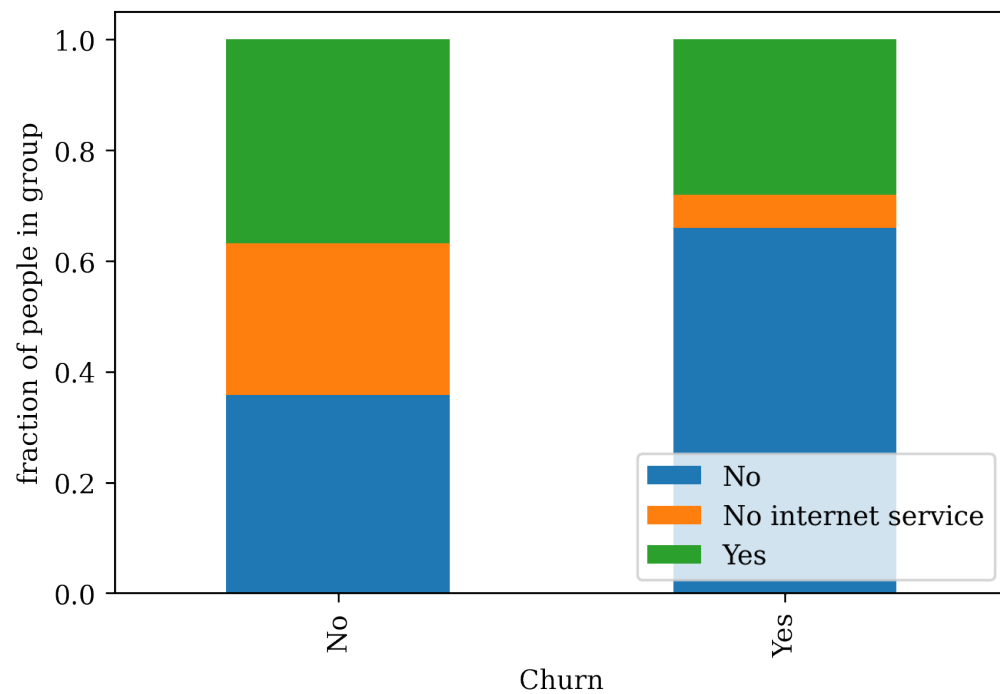
```



```

No          3498
Yes         2019
No internet service  1526
Name: OnlineSecurity, dtype: int64
No          0.496663
Yes         0.286668
No internet service  0.216669
Name: OnlineSecurity, dtype: float64


```



```

No          3088
Yes         2429
No internet service  1526
Name: OnlineBackup, dtype: int64
No          0.438450
Yes         0.344881
No internet service  0.216669
Name: OnlineBackup, dtype: float64


```



```

No          3095
Yes         2422
No internet service  1526
Name: DeviceProtection, dtype: int64
No          0.439443
Yes         0.343888
No internet service  0.216669
Name: DeviceProtection, dtype: float64


```



```

No          3473
Yes         2044
No internet service  1526
Name: TechSupport, dtype: int64
No          0.493114
Yes         0.290217
No internet service  0.216669
Name: TechSupport, dtype: float64


```



```

No                2810
Yes               2707
No internet service  1526
Name: StreamingTV, dtype: int64
No                0.398978
Yes               0.384353
No internet service  0.216669
Name: StreamingTV, dtype: float64


```



```

No                2785
Yes               2732
No internet service 1526
Name: StreamingMovies, dtype: int64
No                0.395428
Yes               0.387903
No internet service 0.216669
Name: StreamingMovies, dtype: float64

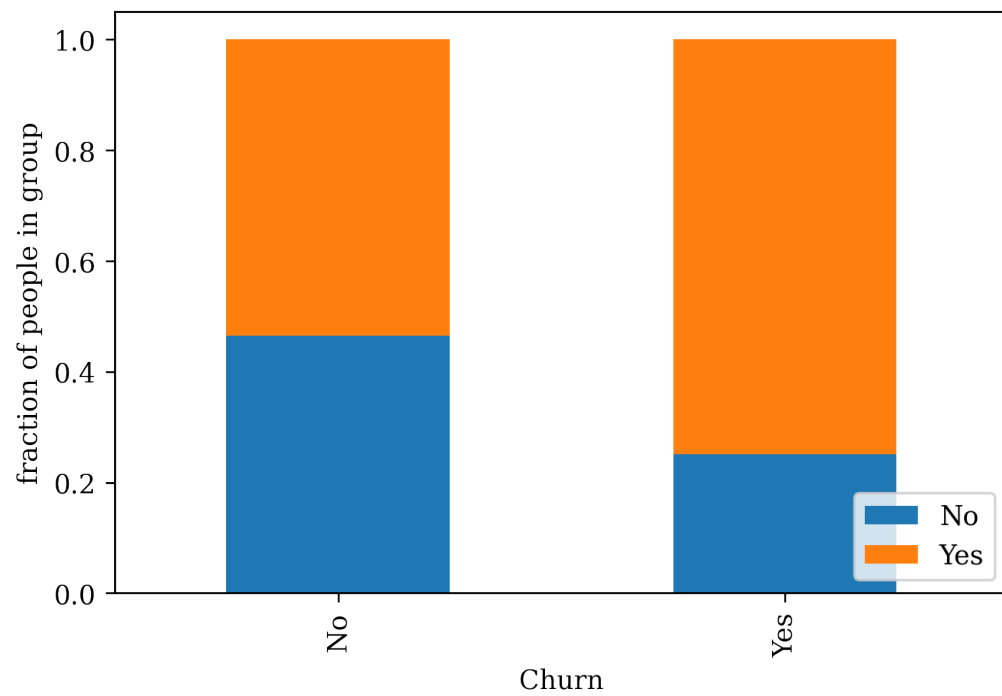
```

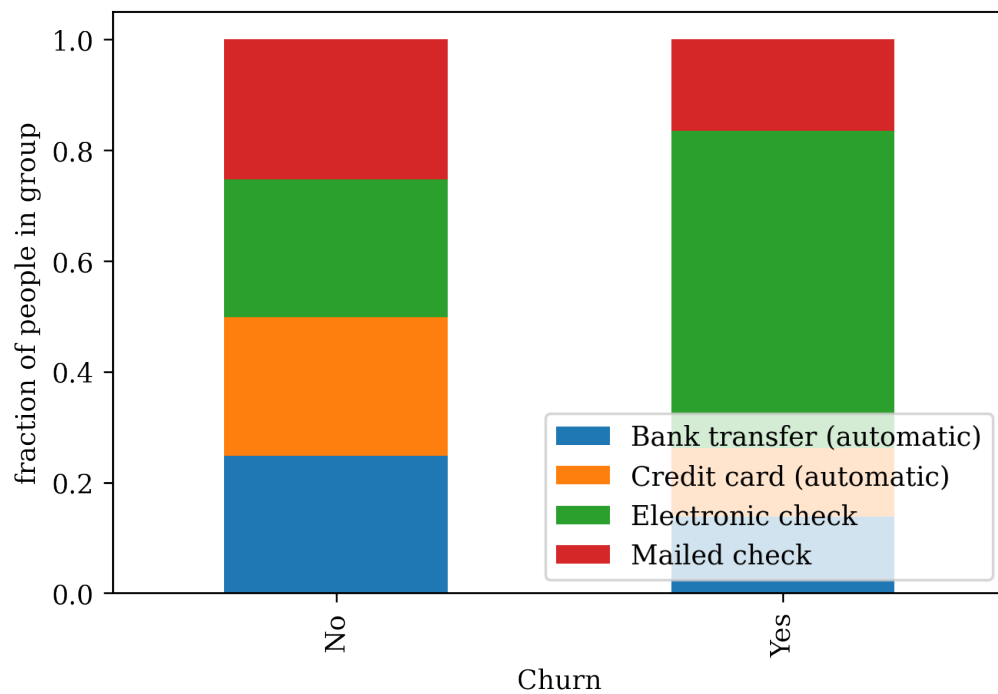
```

Month-to-month    3875
Two year          1695
One year          1473
Name: Contract, dtype: int64
Month-to-month    0.550192
Two year          0.240664
One year          0.209144
Name: Contract, dtype: float64

```




```
Yes    4171
No     2872
Name: PaperlessBilling, dtype: int64
Yes    0.592219
No     0.407781
Name: PaperlessBilling, dtype: float64
```





```

Electronic check      2365
Mailed check         1612
Bank transfer (automatic) 1544
Credit card (automatic) 1522
Name: PaymentMethod, dtype: int64
Electronic check      0.335794
Mailed check         0.228880
Bank transfer (automatic) 0.219225
Credit card (automatic) 0.216101
Name: PaymentMethod, dtype: float64

```




```
count    7043.000000
mean      64.761692
std       30.090047
min       18.250000
25%       35.500000
50%       70.350000
75%       89.850000
max       118.750000
Name: MonthlyCharges, dtype: float64
```



```
count    7032.000000
mean     2283.300441
```

```
std      2266.771362
min       18.800000
25%      401.450000
50%     1397.475000
75%     3794.737500
max     8684.800000
Name: TotalCharges, dtype: float64
```




```
[931]: # Only three continuous variables for scatter matrix

df_continuous = df[['tenure', 'MonthlyCharges', 'TotalCharges']]

pd.plotting.scatter_matrix(df_continuous, figsize=(9, 9),
    ↪marker='o', hist_kws={'bins': 50},
                           s=30, alpha=.1)

plt.show()
```



```
[25]: # Missing Values
print(df.isnull().sum())
```

```
gender          0
SeniorCitizen   0
Partner         0
Dependents      0
tenure          0
PhoneService    0
MultipleLines    0
InternetService  0
OnlineSecurity  0
OnlineBackup    0
DeviceProtection 0
```

```
TechSupport      0
StreamingTV      0
StreamingMovies  0
Contract         0
PaperlessBilling 0
PaymentMethod    0
MonthlyCharges   0
TotalCharges     11
Churn            0
dtype: int64
```

2 2. Methods

```
[30]: # Missing Data: delete 7 rows of missing total charges variable
```

```
df_r = df.dropna()
print(df_r.isnull().sum())
```

```
gender          0
SeniorCitizen   0
Partner         0
Dependents      0
tenure          0
PhoneService    0
MultipleLines   0
InternetService 0
OnlineSecurity  0
OnlineBackup    0
DeviceProtection 0
TechSupport     0
StreamingTV     0
StreamingMovies 0
Contract        0
PaperlessBilling 0
PaymentMethod   0
MonthlyCharges  0
TotalCharges    0
Churn           0
dtype: int64
```

```
[10]: print(df_r.shape[0])
```

```
7032
```

```
[31]: y = df_r['Churn']
X = df_r.loc[:, df_r.columns != 'Churn']
```


2.1 2.1 Evaluation Metric - Baseline F1 Score

```
[12]: # Calculate Baseline f1 Score
      # Predict all classes to be 1

      class_0 = df_r['Churn'].value_counts()[0]
      class_1 = df_r['Churn'].value_counts()[1]

      TN = 0
      FP = class_0

      print(class_0)

      FN = 0
      TP = class_1

      print(class_1)

      p = TP/(TP+FP)
      r = TP/(TP+FN)

      baseline_f1 = (2*p*r)/(p+r)

      print("Hand-calculated baseline F1 score: ",baseline_f1)
```

5163

1869

Hand-calculated baseline F1 score: 0.41995281429052916

```
[13]: y = df_r['Churn'].map(dict(Yes=1, No=0))
      y_baseline = np.full((7032, 1), 1)

      from sklearn.metrics import f1_score
      f1_score(y, y_baseline)

      print("sklearn calculated baseline F1 score: ",baseline_f1)
```

sklearn calculated baseline F1 score: 0.41995281429052916

```
[32]: y = df_r['Churn'].map(dict(Yes=1, No=0))
      y
```

```
[32]: 0      0
      1      0
      2      1
      3      0
      4      1
      ..
```

```

7038    0
7039    0
7040    0
7041    1
7042    0
Name: Churn, Length: 7032, dtype: int64

```

2.2 Machine Learning Models Pipeline

```

[34]: from sklearn.pipeline import make_pipeline
      from sklearn.metrics import mean_squared_error
      from math import sqrt
      from sklearn.model_selection import GridSearchCV
      from sklearn.pipeline import make_pipeline
      from sklearn.compose import ColumnTransformer
      from sklearn.pipeline import Pipeline
      from sklearn.preprocessing import StandardScaler, OneHotEncoder, \
      ↪OrdinalEncoder, LabelEncoder
      from sklearn.model_selection import train_test_split
      from sklearn.metrics import f1_score
      from sklearn.metrics import accuracy_score
      from sklearn.metrics import precision_score
      from sklearn.metrics import recall_score
      from sklearn.model_selection import ParameterGrid

[33]: ordinal_ftrs = \
      ↪['PhoneService', 'MultipleLines', 'InternetService', 'OnlineSecurity', \
      ↪    'OnlineBackup', 'DeviceProtection', 'TechSupport', \
      ↪    'StreamingTV', 'StreamingMovies', 'Contract']
      ordinal_cats = [['No', 'Yes'], ['No phone service', 'No', 'Yes'], ['No', 'DSL', 'Fiber \
      ↪optic'], ['No internet service', 'No', 'Yes'], \
      ↪    ['No internet service', 'No', 'Yes'], ['No internet \
      ↪service', 'No', 'Yes'], ['No internet service', 'No', 'Yes'], \
      ↪    ['No internet service', 'No', 'Yes'], ['No internet \
      ↪service', 'No', 'Yes'], ['Month-to-month', 'One year', 'Two year']]

      cat_ftrs = ['gender', 'SeniorCitizen', 'Partner', \
      ↪    'Dependents', 'PaperlessBilling', 'PaymentMethod']

      num_ftrs = ['tenure', 'MonthlyCharges', 'TotalCharges']

      categorical_transformer = Pipeline(steps=[
          ('onehot', OneHotEncoder(sparse=False, handle_unknown='ignore'))])

```

```

# ordinal encoder
# We need to replace the NaN with a string first!
ordinal_transformer = Pipeline(steps=[
    ('ordinal', OrdinalEncoder(categories = ordinal_cats))])

# standard scaler
numeric_transformer = Pipeline(steps=[
    ('scaler', StandardScaler())])

# collect the transformers
preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, num_ftrs),
        ('cat', categorical_transformer, cat_ftrs),
        ('ord', ordinal_transformer, ordinal_ftrs)])

```

```
[76]: # prep = Pipeline(steps=[('preprocessor', preprocessor)])
```

```
[19]: def MLpipe_Stratify_f1(X, y, preprocessor, ML_algo, param_grid):
    '''
    This function stratified-splits the data to training/validation/test (60/20/
    →20)
    The f1 score as metric score
    '''

    nr_states = 10
    test_scores = np.zeros(nr_states)
    # recall_scores = np.zeros(nr_states)

    val_best_scores = np.zeros(nr_states)
    final_models = []
    feature_importances = np.zeros(nr_states)
    predictions = []

    for i in range(nr_states):

        print('\nrandoms state '+str(i+1))

        X_other, X_test, y_other, y_test = train_test_split(X,y,test_size = 0.
        →2,stratify = y, random_state=22*i)
        X_train, X_val, y_train, y_val =
        →train_test_split(X_other,y_other,test_size = 0.25, stratify = y_other,
        →random_state=22*i)

        train_score = np.zeros(len(ParameterGrid(param_grid)))
        val_score = np.zeros(len(ParameterGrid(param_grid)))

```

```

X_train_prep = preprocessor.fit_transform(X_train)

#         feature_names = preprocessor.transformers_[0][-1] + \
#         list(preprocessor.named_transformers_['cat'][0].
→get_feature_names(cat_fts)) + \
#         preprocessor.transformers_[2][-1]
#         print(feature_names)

X_val_prep = preprocessor.transform(X_val)
X_test_prep = preprocessor.transform(X_test)

models = []

for p in range(len(ParameterGrid(param_grid))):
    params = ParameterGrid(param_grid)[p]
    # print('    ',params)
    try:
        ML = ML_algo(random_state = 22*i)
    except:
        ML = ML_algo()

    ML.set_params(**params)
    ML.fit(X_train_prep,y_train)

    train_score[p] = f1_score(y_train, ML.predict(X_train_prep))

    models.append(ML)
    y_CV_pred = ML.predict(X_val_prep)
    val_score[p]= f1_score(y_val, y_CV_pred)

    # print('    ', 'train score:',train_score[p], 'validation score:
→',val_score[p])

    print([np.argmax(val_score)])
    print(models[np.argmax(val_score)])
    val_best_scores[i] = np.max(val_score)
    print('\nbest model parameters:',ParameterGrid(param_grid)[np.
→argmax(val_score)])
    print('corresponding validation F1 score:',np.max(val_score))

    final_models.append(models[np.argmax(val_score)])

    y_test_pred = final_models[-1].predict(X_test_prep)

    test_scores[i] = f1_score(y_test, y_test_pred)

```

```

    ### include recall_scores as needed

    # recall_scores[i] = recall_score(y_test, y_test_pred)

    print('test F1 score:', test_scores[i])

    return val_best_scores, test_scores, final_models, X_test_prep, y_test

```

2.3 Machine Learning Models Hyperparameters Tuning

```

[327]: # Logistic Regression L1 Regularization

from sklearn.linear_model import LogisticRegression

params = { 'penalty' : ['l1'],
           'C' : [1e-4, 1e-3, 1e-2, 1e-1, 1e0, 1e1, 1e2, 1e3, 1e4],
           'max_iter': [100000],
           'solver': ['saga'] }

Logl1_val_best_F1, Logl1_test_F1, Logl1_models = MLpipe_Stratify_f1(X, y, prep,
↳ LogisticRegression, params)

```

```

randoms state 1
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.0001}
train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.001}
train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.01}
train score: 0.502247191011236 validation score: 0.521172638436482
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.1}
train score: 0.5887016848364717 validation score: 0.5838150289017341
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 1.0}
train score: 0.6076071256620125 validation score: 0.5889046941678521
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 10.0}
train score: 0.6062650602409639 validation score: 0.5957446808510639
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 100.0}
train score: 0.6062650602409639 validation score: 0.594900849858357
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 1000.0}
train score: 0.6062650602409639 validation score: 0.594900849858357
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 10000.0}
train score: 0.6062650602409639 validation score: 0.594900849858357

[5]
LogisticRegression(C=10.0, max_iter=100000, penalty='l1', random_state=0,
                    solver='saga')

```

best model parameters: {'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 10.0}

corresponding validation F1 score: 0.5957446808510639

test F1 score: 0.6169590643274854

randoms state 2

```
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.0001}
train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.001}
train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.01}
train score: 0.5105672969966629 validation score: 0.4966442953020133
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.1}
train score: 0.5967342899554676 validation score: 0.5727272727272728
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 1.0}
train score: 0.6091173617846751 validation score: 0.573134328358209
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 10.0}
train score: 0.6084425036390102 validation score: 0.5819793205317577
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 100.0}
train score: 0.6077669902912621 validation score: 0.5819793205317577
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 1000.0}
train score: 0.6074721009218826 validation score: 0.5819793205317577
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 10000.0}
train score: 0.6074721009218826 validation score: 0.5819793205317577
```

[5]

```
LogisticRegression(C=10.0, max_iter=100000, penalty='l1', random_state=22,
                    solver='saga')
```

best model parameters: {'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 10.0}

corresponding validation F1 score: 0.5819793205317577

test F1 score: 0.5889387144992526

randoms state 3

```
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.0001}
train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.001}
train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.01}
train score: 0.49577464788732395 validation score: 0.5083056478405316
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.1}
train score: 0.5832502492522432 validation score: 0.6005747126436781
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 1.0}
train score: 0.591796875 validation score: 0.5895953757225434
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 10.0}
train score: 0.5941463414634146 validation score: 0.5902578796561604
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 100.0}
train score: 0.5938566552901025 validation score: 0.5894134477825466
```

```

{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 1000.0}
train score: 0.5938566552901025 validation score: 0.5894134477825466
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 10000.0}
train score: 0.5938566552901025 validation score: 0.5894134477825466
[3]
LogisticRegression(C=0.1, max_iter=100000, penalty='l1', random_state=44,
                    solver='saga')

best model parameters: {'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000,
                        'C': 0.1}
corresponding validation F1 score: 0.6005747126436781
test F1 score: 0.5916795069337443

randoms state 4
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.0001}
train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.001}
train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.01}
train score: 0.5172413793103449 validation score: 0.5141955835962145
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.1}
train score: 0.6015779092702169 validation score: 0.5839210155148097
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 1.0}
train score: 0.6069364161849711 validation score: 0.5885634588563459
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 10.0}
train score: 0.6094049904030711 validation score: 0.5905292479108636
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 100.0}
train score: 0.6097794822627037 validation score: 0.592489568845619
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 1000.0}
train score: 0.6097794822627037 validation score: 0.592489568845619
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 10000.0}
train score: 0.6097794822627037 validation score: 0.592489568845619
[6]
LogisticRegression(C=100.0, max_iter=100000, penalty='l1', random_state=66,
                    solver='saga')

best model parameters: {'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000,
                        'C': 100.0}
corresponding validation F1 score: 0.592489568845619
test F1 score: 0.5915080527086385

randoms state 5
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.0001}
train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.001}
train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.01}
train score: 0.5107794361525705 validation score: 0.49423393739703464

```

```

{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.1}
train score: 0.5836228287841191 validation score: 0.5739910313901346
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 1.0}
train score: 0.6081474296799224 validation score: 0.5777777777777777
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 10.0}
train score: 0.6099565007249879 validation score: 0.5823529411764706
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 100.0}
train score: 0.6102514506769826 validation score: 0.5832106038291606
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 1000.0}
train score: 0.6105466860183842 validation score: 0.5852941176470589
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 10000.0}
train score: 0.6105466860183842 validation score: 0.5852941176470589
[7]
LogisticRegression(C=1000.0, max_iter=100000, penalty='l1', random_state=88,
                    solver='saga')

best model parameters: {'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000,
'C': 1000.0}
corresponding validation F1 score: 0.5852941176470589
test F1 score: 0.5965417867435159

randoms state 6
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.0001}
train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.001}
train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.01}
train score: 0.4991587212563096 validation score: 0.5275459098497496
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.1}
train score: 0.5946745562130178 validation score: 0.587183308494784
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 1.0}
train score: 0.6070565490575157 validation score: 0.5947521865889213
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 10.0}
train score: 0.6052123552123552 validation score: 0.6037735849056604
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 100.0}
train score: 0.6061776061776061 validation score: 0.6037735849056604
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 1000.0}
train score: 0.6064703042008691 validation score: 0.6037735849056604
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 10000.0}
train score: 0.6064703042008691 validation score: 0.6037735849056604
[5]
LogisticRegression(C=10.0, max_iter=100000, penalty='l1', random_state=110,
                    solver='saga')

best model parameters: {'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000,
'C': 10.0}
corresponding validation F1 score: 0.6037735849056604
test F1 score: 0.6020260492040521

```


randoms state 7

```
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.0001}
train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.001}
train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.01}
train score: 0.5146005509641873 validation score: 0.48911222780569524
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.1}
train score: 0.5853174603174602 validation score: 0.5735963581183612
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 1.0}
train score: 0.6064703042008691 validation score: 0.5877061469265367
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 10.0}
train score: 0.6081927710843373 validation score: 0.5898203592814371
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 100.0}
train score: 0.6098265895953757 validation score: 0.5898203592814371
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 1000.0}
train score: 0.6091566265060242 validation score: 0.5898203592814371
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 10000.0}
train score: 0.6091566265060242 validation score: 0.5898203592814371
```

[5]

```
LogisticRegression(C=10.0, max_iter=100000, penalty='l1', random_state=132,
                    solver='saga')
```

best model parameters: {'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 10.0}

corresponding validation F1 score: 0.5898203592814371

test F1 score: 0.5982658959537572

randoms state 8

```
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.0001}
train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.001}
train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.01}
train score: 0.5245720596355604 validation score: 0.48344370860927144
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.1}
train score: 0.6164451009354996 validation score: 0.5500747384155457
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 1.0}
train score: 0.6156843643448612 validation score: 0.5755813953488371
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 10.0}
train score: 0.6162847391516334 validation score: 0.5797101449275363
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 100.0}
train score: 0.6153846153846154 validation score: 0.5797101449275363
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 1000.0}
train score: 0.6160583941605839 validation score: 0.5838150289017341
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 10000.0}
train score: 0.6160583941605839 validation score: 0.5838150289017341
```

[7]

```
LogisticRegression(C=1000.0, max_iter=100000, penalty='l1', random_state=154,  
                    solver='saga')
```

best model parameters: {'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000,
'C': 1000.0}

corresponding validation F1 score: 0.5838150289017341

test F1 score: 0.5861561119293078

randoms state 9

```
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.0001}  
train score: 0.0 validation score: 0.0  
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.001}  
train score: 0.0 validation score: 0.0  
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.01}  
train score: 0.4991587212563096 validation score: 0.4916387959866221  
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.1}  
train score: 0.5879446640316206 validation score: 0.587887740029542  
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 1.0}  
train score: 0.5988372093023256 validation score: 0.6034985422740524  
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 10.0}  
train score: 0.6004842615012106 validation score: 0.6075581395348837  
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 100.0}  
train score: 0.6005802707930368 validation score: 0.6055312954876273  
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 1000.0}  
train score: 0.5999032414126754 validation score: 0.6055312954876273  
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 10000.0}  
train score: 0.5999032414126754 validation score: 0.6055312954876273
```

[5]

```
LogisticRegression(C=10.0, max_iter=100000, penalty='l1', random_state=176,  
                    solver='saga')
```

best model parameters: {'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000,
'C': 10.0}

corresponding validation F1 score: 0.6075581395348837

test F1 score: 0.5920471281296025

randoms state 10

```
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.0001}  
train score: 0.0 validation score: 0.0  
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.001}  
train score: 0.0 validation score: 0.0  
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.01}  
train score: 0.5176211453744494 validation score: 0.4983498349834984  
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 0.1}  
train score: 0.5954500494559841 validation score: 0.5697329376854601  
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 1.0}  
train score: 0.6085271317829458 validation score: 0.5840455840455839
```

```
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 10.0}
train score: 0.6070565490575157 validation score: 0.5820256776034237
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 100.0}
train score: 0.6070565490575157 validation score: 0.5820256776034237
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 1000.0}
train score: 0.6070565490575157 validation score: 0.5820256776034237
{'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000, 'C': 10000.0}
train score: 0.6070565490575157 validation score: 0.5820256776034237
```

[4]

```
LogisticRegression(max_iter=100000, penalty='l1', random_state=198,
                    solver='saga')
```

```
best model parameters: {'solver': 'saga', 'penalty': 'l1', 'max_iter': 100000,
                        'C': 1.0}
```

```
corresponding validation F1 score: 0.5840455840455839
```

```
test F1 score: 0.6
```

```
[868]: # Saving Logistic L1 Regularization Best Models at 10 Random States
filename = direct + '/results/log_l1_best_models.sav'
joblib.dump(Logl1_models, filename)
```

```
[868]: ['/Users/liyuetian1/Documents/GitHub/DATA1030_MidtermProject/results/log_l1_best
_models.sav']
```

```
[21]: # Logistic Regression L2
from sklearn.linear_model import LogisticRegression
params = { 'penalty' : ['l2'],
           'C' : [1e-4, 1e-3, 1e-2, 1e-1, 1e0, 1e1, 1e2, 1e3, 1e4],
           'max_iter': [10000],
           'solver': ['saga'] }

Logl2_val_best_F1, Logl2_test_F1, Logl2_models, X_test, Y_test = MLpipe_Stratify_f1(X, y, preprocessor, LogisticRegression, params)

# To include recall score:
# Logl2_val_best_F1, Logl2_test_F1, Logl2_models, X_test, Y_test, recall_scores = MLpipe_Stratify_f1(X, y, preprocessor, LogisticRegression, params)
```

```
randoms state 1
```

[5]

```
LogisticRegression(C=10.0, max_iter=10000, random_state=0, solver='saga')
```

```
best model parameters: {'solver': 'saga', 'penalty': 'l2', 'max_iter': 10000,
                        'C': 10.0}
```

```
corresponding validation F1 score: 0.5957446808510639
```

```
test F1 score: 0.6169590643274854
```

```

randoms state 2
[5]
LogisticRegression(C=10.0, max_iter=10000, random_state=22, solver='saga')

best model parameters: {'solver': 'saga', 'penalty': 'l2', 'max_iter': 10000,
'C': 10.0}
corresponding validation F1 score: 0.5819793205317577
test F1 score: 0.5889387144992526

randoms state 3
[3]
LogisticRegression(C=0.1, max_iter=10000, random_state=44, solver='saga')

best model parameters: {'solver': 'saga', 'penalty': 'l2', 'max_iter': 10000,
'C': 0.1}
corresponding validation F1 score: 0.5936599423631125
test F1 score: 0.5969230769230769

randoms state 4
[6]
LogisticRegression(C=100.0, max_iter=10000, random_state=66, solver='saga')

best model parameters: {'solver': 'saga', 'penalty': 'l2', 'max_iter': 10000,
'C': 100.0}
corresponding validation F1 score: 0.592489568845619
test F1 score: 0.5915080527086385

randoms state 5
[6]
LogisticRegression(C=100.0, max_iter=10000, random_state=88, solver='saga')

best model parameters: {'solver': 'saga', 'penalty': 'l2', 'max_iter': 10000,
'C': 100.0}
corresponding validation F1 score: 0.5852941176470589
test F1 score: 0.5965417867435159

randoms state 6
[5]
LogisticRegression(C=10.0, max_iter=10000, random_state=110, solver='saga')

best model parameters: {'solver': 'saga', 'penalty': 'l2', 'max_iter': 10000,
'C': 10.0}
corresponding validation F1 score: 0.6037735849056604
test F1 score: 0.6

randoms state 7
[4]

```

```
LogisticRegression(max_iter=10000, random_state=132, solver='saga')

best model parameters: {'solver': 'saga', 'penalty': 'l2', 'max_iter': 10000,
'C': 1.0}
corresponding validation F1 score: 0.5937031484257871
test F1 score: 0.5982658959537572
```

```
randoms state 8
[7]
LogisticRegression(C=1000.0, max_iter=10000, random_state=154, solver='saga')

best model parameters: {'solver': 'saga', 'penalty': 'l2', 'max_iter': 10000,
'C': 1000.0}
corresponding validation F1 score: 0.5838150289017341
test F1 score: 0.5861561119293078
```

```
randoms state 9
[4]
LogisticRegression(max_iter=10000, random_state=176, solver='saga')

best model parameters: {'solver': 'saga', 'penalty': 'l2', 'max_iter': 10000,
'C': 1.0}
corresponding validation F1 score: 0.6075581395348837
test F1 score: 0.5941176470588235
```

```
randoms state 10
[5]
LogisticRegression(C=10.0, max_iter=10000, random_state=198, solver='saga')

best model parameters: {'solver': 'saga', 'penalty': 'l2', 'max_iter': 10000,
'C': 10.0}
corresponding validation F1 score: 0.582857142857143
test F1 score: 0.5988372093023255
```

```
[22]: np.mean(recall_scores)
```

```
[22]: 0.5433155080213903
```

```
[23]: np.std(recall_scores)
```

```
[23]: 0.013517887336646571
```

```
[871]: # Saving Logistic L2 Regularization Best Models at 10 Random States

filename = direct + '/results/log_l2_best_models.sav'
joblib.dump(Logl2_models, filename)
```

```
[871]: ['/Users/liyuetian1/Documents/GitHub/DATA1030_MidtermProject/results/log_l2_best_models.sav']
```

```
[323]: # Logistic Regression Elastic Net
```

```
params = { 'penalty' : ['elasticnet'],  
           'C' : [1e-4, 1e-3, 1e-2, 1e-1, 1e0, 1e1, 1e2, 1e3, 1e4],  
           'l1_ratio': [0.01, 0.1, 0.25, 0.5, 0.75, 0.9, 0.99],  
           'max_iter': [100000],  
           'solver': ['saga'] }  
  
LogEN_val_best_F1, LogEN_test_F1, LogEN_models = MLpipe_Stratify_f1(X, y,  
→preprocessor, LogisticRegression, params)
```

```
randoms state 1  
  {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':  
0.01, 'C': 0.0001}  
  train score: 0.0 validation score: 0.0  
  {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':  
0.1, 'C': 0.0001}  
  train score: 0.0 validation score: 0.0  
  {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':  
0.25, 'C': 0.0001}  
  train score: 0.0 validation score: 0.0  
  {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':  
0.5, 'C': 0.0001}  
  train score: 0.0 validation score: 0.0  
  {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':  
0.75, 'C': 0.0001}  
  train score: 0.0 validation score: 0.0  
  {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':  
0.9, 'C': 0.0001}  
  train score: 0.0 validation score: 0.0  
  {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':  
0.99, 'C': 0.0001}  
  train score: 0.0 validation score: 0.0  
  {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':  
0.01, 'C': 0.001}  
  train score: 0.18354430379746836 validation score: 0.2186046511627907  
  {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':  
0.1, 'C': 0.001}  
  train score: 0.0 validation score: 0.0  
  {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':  
0.25, 'C': 0.001}  
  train score: 0.0 validation score: 0.0  
  {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
```

```

0.5, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 0.01}
    train score: 0.5609504132231405 validation score: 0.5606060606060607
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 0.01}
    train score: 0.5601249349297241 validation score: 0.5540334855403348
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 0.01}
    train score: 0.5487480021310602 validation score: 0.5412130637636081
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 0.01}
    train score: 0.5277015907844212 validation score: 0.5278219395866455
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 0.01}
    train score: 0.5139353400222966 validation score: 0.5283630470016207
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 0.01}
    train score: 0.5061728395061729 validation score: 0.5252854812398042
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 0.01}
    train score: 0.502247191011236 validation score: 0.521172638436482
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 0.1}
    train score: 0.6018563751831949 validation score: 0.5853658536585366
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 0.1}
    train score: 0.599706026457619 validation score: 0.5865522174535049
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 0.1}
    train score: 0.6013712047012733 validation score: 0.5882352941176471
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 0.1}
    train score: 0.6018654884634266 validation score: 0.5899280575539568
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 0.1}
    train score: 0.594381468703795 validation score: 0.5846599131693199
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':

```

```

0.9, 'C': 0.1}
    train score: 0.5898070262246413 validation score: 0.5846599131693199
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 0.1}
    train score: 0.5887016848364717 validation score: 0.5838150289017341
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 1.0}
    train score: 0.6065573770491803 validation score: 0.59375
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 1.0}
    train score: 0.6062650602409639 validation score: 0.59375
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 1.0}
    train score: 0.6062650602409639 validation score: 0.5909090909090908
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 1.0}
    train score: 0.6076071256620125 validation score: 0.5909090909090908
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 1.0}
    train score: 0.6076071256620125 validation score: 0.5909090909090908
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 1.0}
    train score: 0.6085700529610014 validation score: 0.5909090909090908
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 1.0}
    train score: 0.6076071256620125 validation score: 0.5889046941678521
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 10.0}
    train score: 0.6062650602409639 validation score: 0.5957446808510639
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 10.0}
    train score: 0.6062650602409639 validation score: 0.5957446808510639
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 10.0}
    train score: 0.6062650602409639 validation score: 0.5957446808510639
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 10.0}
    train score: 0.6062650602409639 validation score: 0.5957446808510639
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 10.0}
    train score: 0.6062650602409639 validation score: 0.5957446808510639
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 10.0}
    train score: 0.6062650602409639 validation score: 0.5957446808510639
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 10.0}
    train score: 0.6062650602409639 validation score: 0.5957446808510639
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':

```



```

0.25, 'C': 10000.0}
    train score: 0.6062650602409639 validation score: 0.594900849858357
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 10000.0}
    train score: 0.6062650602409639 validation score: 0.594900849858357
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 10000.0}
    train score: 0.6062650602409639 validation score: 0.594900849858357
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 10000.0}
    train score: 0.6062650602409639 validation score: 0.594900849858357
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 10000.0}
    train score: 0.6062650602409639 validation score: 0.594900849858357
[35]
LogisticRegression(C=10.0, l1_ratio=0.01, max_iter=100000, penalty='elasticnet',
                    random_state=0, solver='saga')

```

```

best model parameters: {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter':
100000, 'l1_ratio': 0.01, 'C': 10.0}
corresponding validation F1 score: 0.5957446808510639
test F1 score: 0.6169590643274854

```

```

randoms state 2
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 0.001}
    train score: 0.19370078740157484 validation score: 0.1375921375921376

```

```

    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 0.01}
    train score: 0.5767060030785017 validation score: 0.5385826771653544
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 0.01}
    train score: 0.5689210118740321 validation score: 0.5209003215434084
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 0.01}
    train score: 0.5601265822784809 validation score: 0.5244299674267101
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 0.01}
    train score: 0.53470715835141 validation score: 0.5074626865671642
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 0.01}
    train score: 0.525909592061742 validation score: 0.5041736227045076
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 0.01}
    train score: 0.5153032832498609 validation score: 0.5033333333333333
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 0.01}
    train score: 0.5103294249022893 validation score: 0.49916247906197664
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 0.1}
    train score: 0.6056751467710371 validation score: 0.5684210526315789
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 0.1}
    train score: 0.605288932419197 validation score: 0.5705705705705707
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 0.1}
    train score: 0.6031434184675836 validation score: 0.5748502994011976

```

```

    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 0.1}
    train score: 0.6012814194184327 validation score: 0.5787106446776612
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 0.1}
    train score: 0.5995061728395061 validation score: 0.579185520361991
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 0.1}
    train score: 0.5985185185185184 validation score: 0.5783132530120482
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 0.1}
    train score: 0.597132970835393 validation score: 0.5770392749244713
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 1.0}
    train score: 0.6104651162790697 validation score: 0.5786350148367952
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 1.0}
    train score: 0.6107610276296654 validation score: 0.5765230312035662
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 1.0}
    train score: 0.6107610276296654 validation score: 0.5756676557863502
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 1.0}
    train score: 0.6091173617846751 validation score: 0.5777777777777777
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 1.0}
    train score: 0.6087378640776699 validation score: 0.5786350148367952
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 1.0}
    train score: 0.6091173617846751 validation score: 0.5773809523809523
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 1.0}
    train score: 0.6091173617846751 validation score: 0.573134328358209
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 10.0}
    train score: 0.6084425036390102 validation score: 0.5819793205317577
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 10.0}
    train score: 0.6084425036390102 validation score: 0.5819793205317577
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 10.0}
    train score: 0.6084425036390102 validation score: 0.5819793205317577
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 10.0}
    train score: 0.6084425036390102 validation score: 0.5819793205317577
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 10.0}
    train score: 0.6084425036390102 validation score: 0.5819793205317577

```



```

    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 10000.0}
    train score: 0.6074721009218826 validation score: 0.5819793205317577
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 10000.0}
    train score: 0.6074721009218826 validation score: 0.5819793205317577
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 10000.0}
    train score: 0.6074721009218826 validation score: 0.5819793205317577
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 10000.0}
    train score: 0.6074721009218826 validation score: 0.5819793205317577
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 10000.0}
    train score: 0.6074721009218826 validation score: 0.5819793205317577
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 10000.0}
    train score: 0.6074721009218826 validation score: 0.5819793205317577
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 10000.0}
    train score: 0.6074721009218826 validation score: 0.5819793205317577
[35]

```

```

LogisticRegression(C=10.0, l1_ratio=0.01, max_iter=100000, penalty='elasticnet',
                    random_state=22, solver='saga')

```

```

best model parameters: {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter':
100000, 'l1_ratio': 0.01, 'C': 10.0}
corresponding validation F1 score: 0.5819793205317577
test F1 score: 0.5889387144992526

```

randoms state 3

```

    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 0.0001}

```

```

train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 0.0001}
train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 0.001}
train score: 0.15384615384615385 validation score: 0.11793611793611795
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 0.001}
train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 0.001}
train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 0.001}
train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 0.001}
train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 0.001}
train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 0.001}
train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 0.01}
train score: 0.5507853403141362 validation score: 0.5535168195718654
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 0.01}
train score: 0.5472794506075014 validation score: 0.5432098765432098
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 0.01}
train score: 0.534048257372654 validation score: 0.5339652448657187
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 0.01}
train score: 0.5219298245614036 validation score: 0.5281803542673108
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 0.01}
train score: 0.5102834908282379 validation score: 0.5154975530179445
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 0.01}
train score: 0.5036537380550871 validation score: 0.5114754098360657
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 0.01}
train score: 0.49577464788732395 validation score: 0.5083056478405316
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 0.1}

```

```

train score: 0.5828797624938149 validation score: 0.5936599423631125
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 0.1}
train score: 0.5821782178217821 validation score: 0.5936599423631125
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 0.1}
train score: 0.5829195630585898 validation score: 0.5931232091690544
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 0.1}
train score: 0.5819631290483308 validation score: 0.5974395448079659
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 0.1}
train score: 0.5850746268656716 validation score: 0.592274678111588
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 0.1}
train score: 0.5835411471321695 validation score: 0.599713055954089
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 0.1}
train score: 0.5832502492522432 validation score: 0.6005747126436781
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 1.0}
train score: 0.593460224499756 validation score: 0.5919540229885057
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 1.0}
train score: 0.593460224499756 validation score: 0.5899280575539568
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 1.0}
train score: 0.593460224499756 validation score: 0.5870503597122301
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 1.0}
train score: 0.5930630190522717 validation score: 0.5887445887445887
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 1.0}
train score: 0.5933528836754642 validation score: 0.5895953757225434
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 1.0}
train score: 0.5920859794821691 validation score: 0.5895953757225434
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 1.0}
train score: 0.591796875 validation score: 0.5895953757225434
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 10.0}
train score: 0.5941463414634146 validation score: 0.5894134477825466
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 10.0}
train score: 0.5941463414634146 validation score: 0.5894134477825466
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 10.0}

```



```

train score: 0.5938566552901025 validation score: 0.5894134477825466
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 1000.0}
train score: 0.5938566552901025 validation score: 0.5894134477825466
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 1000.0}
train score: 0.5938566552901025 validation score: 0.5894134477825466
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 10000.0}
train score: 0.5938566552901025 validation score: 0.5894134477825466
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 10000.0}
train score: 0.5938566552901025 validation score: 0.5894134477825466
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 10000.0}
train score: 0.5938566552901025 validation score: 0.5894134477825466
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 10000.0}
train score: 0.5938566552901025 validation score: 0.5894134477825466
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 10000.0}
train score: 0.5938566552901025 validation score: 0.5894134477825466
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 10000.0}
train score: 0.5938566552901025 validation score: 0.5894134477825466
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 10000.0}
train score: 0.5938566552901025 validation score: 0.5894134477825466

```

[27]

```

LogisticRegression(C=0.1, l1_ratio=0.99, max_iter=100000, penalty='elasticnet',
                    random_state=44, solver='saga')

```

```

best model parameters: {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter':
100000, 'l1_ratio': 0.99, 'C': 0.1}

```

```

corresponding validation F1 score: 0.6005747126436781

```

```

test F1 score: 0.5916795069337443

```

randoms state 4

```

{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 0.0001}
train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 0.0001}
train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 0.0001}
train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':

```

```

0.5, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 0.001}
    train score: 0.19607843137254902 validation score: 0.21495327102803738
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 0.01}
    train score: 0.574083634486319 validation score: 0.5542168674698795
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 0.01}
    train score: 0.5689027561102444 validation score: 0.5460030165912519
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 0.01}
    train score: 0.5585585585585585 validation score: 0.5482388973966309
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 0.01}
    train score: 0.5326027397260273 validation score: 0.5360501567398119
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 0.01}
    train score: 0.5193370165745858 validation score: 0.5266457680250783
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':

```

```

0.9, 'C': 0.01}
    train score: 0.5177777777777778 validation score: 0.5220125786163522
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 0.01}
    train score: 0.5172413793103449 validation score: 0.5141955835962145
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 0.1}
    train score: 0.6038104543234002 validation score: 0.5907172995780591
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 0.1}
    train score: 0.6051732552464617 validation score: 0.5875706214689266
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 0.1}
    train score: 0.6029411764705883 validation score: 0.5875706214689266
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 0.1}
    train score: 0.6036256736893679 validation score: 0.5864022662889519
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 0.1}
    train score: 0.6034398034398034 validation score: 0.5867418899858956
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 0.1}
    train score: 0.6034398034398034 validation score: 0.5859154929577466
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 0.1}
    train score: 0.6015779092702169 validation score: 0.5839210155148097
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 1.0}
    train score: 0.6073147256977862 validation score: 0.5885634588563459
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 1.0}
    train score: 0.6073147256977862 validation score: 0.5885634588563459
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 1.0}
    train score: 0.6073147256977862 validation score: 0.5885634588563459
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 1.0}
    train score: 0.6066441983630236 validation score: 0.5885634588563459
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 1.0}
    train score: 0.6072289156626506 validation score: 0.5885634588563459
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 1.0}
    train score: 0.6069364161849711 validation score: 0.5885634588563459
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 1.0}
    train score: 0.6069364161849711 validation score: 0.5885634588563459
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':

```

```

0.01, 'C': 10.0}
    train score: 0.6094049904030711 validation score: 0.5905292479108636
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 10.0}
    train score: 0.6094049904030711 validation score: 0.5905292479108636
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 10.0}
    train score: 0.6100719424460431 validation score: 0.5905292479108636
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 10.0}
    train score: 0.6100719424460431 validation score: 0.5905292479108636
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 10.0}
    train score: 0.6094049904030711 validation score: 0.5905292479108636
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 10.0}
    train score: 0.6094049904030711 validation score: 0.5905292479108636
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 10.0}
    train score: 0.6094049904030711 validation score: 0.5905292479108636
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 100.0}
    train score: 0.6097794822627037 validation score: 0.592489568845619
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 100.0}
    train score: 0.6097794822627037 validation score: 0.592489568845619
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 100.0}
    train score: 0.6097794822627037 validation score: 0.592489568845619
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 100.0}
    train score: 0.6097794822627037 validation score: 0.592489568845619
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 100.0}
    train score: 0.6097794822627037 validation score: 0.592489568845619
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 100.0}
    train score: 0.6097794822627037 validation score: 0.592489568845619
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 100.0}
    train score: 0.6097794822627037 validation score: 0.592489568845619
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 1000.0}
    train score: 0.6097794822627037 validation score: 0.592489568845619
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 1000.0}
    train score: 0.6097794822627037 validation score: 0.592489568845619
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':

```

```

0.25, 'C': 1000.0}
    train score: 0.6097794822627037 validation score: 0.592489568845619
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 1000.0}
    train score: 0.6097794822627037 validation score: 0.592489568845619
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 1000.0}
    train score: 0.6097794822627037 validation score: 0.592489568845619
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 1000.0}
    train score: 0.6097794822627037 validation score: 0.592489568845619
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 1000.0}
    train score: 0.6097794822627037 validation score: 0.592489568845619
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 10000.0}
    train score: 0.6097794822627037 validation score: 0.592489568845619
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 10000.0}
    train score: 0.6097794822627037 validation score: 0.592489568845619
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 10000.0}
    train score: 0.6097794822627037 validation score: 0.592489568845619
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 10000.0}
    train score: 0.6097794822627037 validation score: 0.592489568845619
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 10000.0}
    train score: 0.6097794822627037 validation score: 0.592489568845619
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 10000.0}
    train score: 0.6097794822627037 validation score: 0.592489568845619
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 10000.0}
    train score: 0.6097794822627037 validation score: 0.592489568845619
[42]
LogisticRegression(C=100.0, l1_ratio=0.01, max_iter=100000,
                    penalty='elasticnet', random_state=66, solver='saga')

best model parameters: {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter':
100000, 'l1_ratio': 0.01, 'C': 100.0}
corresponding validation F1 score: 0.592489568845619
test F1 score: 0.5915080527086385

randoms state 5
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 0.0001}
    train score: 0.0 validation score: 0.0

```

```

    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 0.001}
    train score: 0.18383518225039622 validation score: 0.18527315914489312
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 0.01}
    train score: 0.5622739018087856 validation score: 0.5598755832037324
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 0.01}
    train score: 0.5567765567765568 validation score: 0.5460317460317461
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 0.01}
    train score: 0.5458399576046635 validation score: 0.5399361022364216

```

```

    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 0.01}
    train score: 0.5342019543973943 validation score: 0.5194805194805195
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 0.01}
    train score: 0.5168788046485888 validation score: 0.5090311986863709
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 0.01}
    train score: 0.5108273181565798 validation score: 0.5
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 0.01}
    train score: 0.5086254869226489 validation score: 0.49586776859504134
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 0.1}
    train score: 0.5932872655478777 validation score: 0.5769805680119581
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 0.1}
    train score: 0.5925925925925927 validation score: 0.5761194029850746
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 0.1}
    train score: 0.5935802469135801 validation score: 0.5765765765765766
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 0.1}
    train score: 0.5899209486166008 validation score: 0.5748502994011976
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 0.1}
    train score: 0.5893385982230998 validation score: 0.5739910313901346
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 0.1}
    train score: 0.5841584158415841 validation score: 0.5748502994011976
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 0.1}
    train score: 0.5843253968253969 validation score: 0.5739910313901346
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 1.0}
    train score: 0.6067961165048543 validation score: 0.5798816568047337
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 1.0}
    train score: 0.6074721009218826 validation score: 0.5798816568047337
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 1.0}
    train score: 0.6065891472868218 validation score: 0.5798816568047337
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 1.0}
    train score: 0.6078526417838099 validation score: 0.5798816568047337
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 1.0}
    train score: 0.6074721009218826 validation score: 0.5798816568047337

```



```

    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 1.0}
    train score: 0.6081474296799224 validation score: 0.5777777777777777
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 1.0}
    train score: 0.6081474296799224 validation score: 0.5777777777777777
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 10.0}
    train score: 0.6105466860183842 validation score: 0.5832106038291606
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 10.0}
    train score: 0.6105466860183842 validation score: 0.5832106038291606
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 10.0}
    train score: 0.6105466860183842 validation score: 0.5832106038291606
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 10.0}
    train score: 0.6102514506769826 validation score: 0.5823529411764706
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 10.0}
    train score: 0.6102514506769826 validation score: 0.5823529411764706
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 10.0}
    train score: 0.6099565007249879 validation score: 0.5823529411764706
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 10.0}
    train score: 0.6099565007249879 validation score: 0.5823529411764706
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 100.0}
    train score: 0.6105466860183842 validation score: 0.5852941176470589
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 100.0}
    train score: 0.6105466860183842 validation score: 0.5852941176470589
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 100.0}
    train score: 0.6105466860183842 validation score: 0.5852941176470589
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 100.0}
    train score: 0.6105466860183842 validation score: 0.5852941176470589
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 100.0}
    train score: 0.6102514506769826 validation score: 0.5832106038291606
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 100.0}
    train score: 0.6102514506769826 validation score: 0.5832106038291606
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 100.0}
    train score: 0.6102514506769826 validation score: 0.5832106038291606

```

```

    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 1000.0}
    train score: 0.6105466860183842 validation score: 0.5852941176470589
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 1000.0}
    train score: 0.6105466860183842 validation score: 0.5852941176470589
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 1000.0}
    train score: 0.6105466860183842 validation score: 0.5852941176470589
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 1000.0}
    train score: 0.6105466860183842 validation score: 0.5852941176470589
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 1000.0}
    train score: 0.6105466860183842 validation score: 0.5852941176470589
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 1000.0}
    train score: 0.6105466860183842 validation score: 0.5852941176470589
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 1000.0}
    train score: 0.6105466860183842 validation score: 0.5852941176470589
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 10000.0}
    train score: 0.6105466860183842 validation score: 0.5852941176470589
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 10000.0}
    train score: 0.6105466860183842 validation score: 0.5852941176470589
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 10000.0}
    train score: 0.6105466860183842 validation score: 0.5852941176470589
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 10000.0}
    train score: 0.6105466860183842 validation score: 0.5852941176470589
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 10000.0}
    train score: 0.6105466860183842 validation score: 0.5852941176470589
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 10000.0}
    train score: 0.6105466860183842 validation score: 0.5852941176470589
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 10000.0}
    train score: 0.6105466860183842 validation score: 0.5852941176470589
[42]
LogisticRegression(C=100.0, l1_ratio=0.01, max_iter=100000,
                    penalty='elasticnet', random_state=88, solver='saga')

best model parameters: {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter':
100000, 'l1_ratio': 0.01, 'C': 100.0}

```

corresponding validation F1 score: 0.5852941176470589
test F1 score: 0.5965417867435159

randoms state 6

```
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':  
0.01, 'C': 0.0001}  
train score: 0.0 validation score: 0.0  
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':  
0.1, 'C': 0.0001}  
train score: 0.0 validation score: 0.0  
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':  
0.25, 'C': 0.0001}  
train score: 0.0 validation score: 0.0  
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':  
0.5, 'C': 0.0001}  
train score: 0.0 validation score: 0.0  
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':  
0.75, 'C': 0.0001}  
train score: 0.0 validation score: 0.0  
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':  
0.9, 'C': 0.0001}  
train score: 0.0 validation score: 0.0  
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':  
0.99, 'C': 0.0001}  
train score: 0.0 validation score: 0.0  
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':  
0.01, 'C': 0.001}  
train score: 0.16051364365971107 validation score: 0.18138424821002389  
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':  
0.1, 'C': 0.001}  
train score: 0.0 validation score: 0.0  
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':  
0.25, 'C': 0.001}  
train score: 0.0 validation score: 0.0  
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':  
0.5, 'C': 0.001}  
train score: 0.0 validation score: 0.0  
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':  
0.75, 'C': 0.001}  
train score: 0.0 validation score: 0.0  
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':  
0.9, 'C': 0.001}  
train score: 0.0 validation score: 0.0  
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':  
0.99, 'C': 0.001}  
train score: 0.0 validation score: 0.0  
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':  
0.01, 'C': 0.01}
```

```

train score: 0.5621454357916451 validation score: 0.571875
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 0.01}
train score: 0.5580426861009891 validation score: 0.5687203791469194
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 0.01}
train score: 0.5428265524625266 validation score: 0.5654952076677315
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 0.01}
train score: 0.5206384149697303 validation score: 0.5415986949429037
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 0.01}
train score: 0.50917176209005 validation score: 0.5250836120401338
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 0.01}
train score: 0.5047459519821329 validation score: 0.53
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 0.01}
train score: 0.5033482142857143 validation score: 0.5291181364392677
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 0.1}
train score: 0.5965601965601965 validation score: 0.5891016200294551
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 0.1}
train score: 0.5975429975429974 validation score: 0.5941176470588235
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 0.1}
train score: 0.5969563082965145 validation score: 0.591715976331361
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 0.1}
train score: 0.5981308411214954 validation score: 0.5929203539823009
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 0.1}
train score: 0.5973385904386398 validation score: 0.5829596412556053
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 0.1}
train score: 0.5936883629191321 validation score: 0.5868263473053892
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 0.1}
train score: 0.5936883629191321 validation score: 0.587183308494784
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 1.0}
train score: 0.6077294685990339 validation score: 0.5997088791848617
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 1.0}
train score: 0.6077294685990339 validation score: 0.5997088791848617
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 1.0}

```

```

train score: 0.6077294685990339 validation score: 0.5997088791848617
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 1.0}
train score: 0.6070565490575157 validation score: 0.5976676384839649
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 1.0}
train score: 0.6073500967117988 validation score: 0.5967976710334788
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 1.0}
train score: 0.6070565490575157 validation score: 0.5947521865889213
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 1.0}
train score: 0.6070565490575157 validation score: 0.5947521865889213
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 10.0}
train score: 0.6048309178743961 validation score: 0.6037735849056604
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 10.0}
train score: 0.6051232479458675 validation score: 0.6037735849056604
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 10.0}
train score: 0.6051232479458675 validation score: 0.6037735849056604
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 10.0}
train score: 0.6048309178743961 validation score: 0.6037735849056604
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 10.0}
train score: 0.6052123552123552 validation score: 0.6037735849056604
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 10.0}
train score: 0.6052123552123552 validation score: 0.6037735849056604
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 10.0}
train score: 0.6052123552123552 validation score: 0.6037735849056604
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 100.0}
train score: 0.6057971014492753 validation score: 0.6037735849056604
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 100.0}
train score: 0.6057971014492753 validation score: 0.6037735849056604
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 100.0}
train score: 0.6057971014492753 validation score: 0.6037735849056604
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 100.0}
train score: 0.6057971014492753 validation score: 0.6037735849056604
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 100.0}

```



```

train score: 0.6064703042008691 validation score: 0.6037735849056604
[35]
LogisticRegression(C=10.0, l1_ratio=0.01, max_iter=100000, penalty='elasticnet',
                    random_state=110, solver='saga')

best model parameters: {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter':
100000, 'l1_ratio': 0.01, 'C': 10.0}
corresponding validation F1 score: 0.6037735849056604
test F1 score: 0.6

randoms state 7
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 0.001}
    train score: 0.18354430379746836 validation score: 0.1686746987951807
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':

```

```

0.9, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 0.01}
    train score: 0.5633074935400516 validation score: 0.5468998410174881
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 0.01}
    train score: 0.5617860851505712 validation score: 0.5408
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 0.01}
    train score: 0.5514316012725344 validation score: 0.5492730210016156
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 0.01}
    train score: 0.5345572354211662 validation score: 0.506578947368421
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 0.01}
    train score: 0.5289617486338798 validation score: 0.4966887417218543
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 0.01}
    train score: 0.520065970313359 validation score: 0.4983388704318937
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 0.01}
    train score: 0.5146005509641873 validation score: 0.48911222780569524
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 0.1}
    train score: 0.5970588235294118 validation score: 0.592814371257485
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 0.1}
    train score: 0.5973516429622364 validation score: 0.592814371257485
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 0.1}
    train score: 0.5984251968503937 validation score: 0.5873493975903614
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 0.1}
    train score: 0.5956607495069034 validation score: 0.5873493975903614
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 0.1}
    train score: 0.592482690405539 validation score: 0.5864661654135338
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 0.1}
    train score: 0.5902125556104796 validation score: 0.5735963581183612
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 0.1}
    train score: 0.5856079404466501 validation score: 0.5714285714285714
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':

```



```

0.01, 'C': 1.0}
    train score: 0.6061776061776061 validation score: 0.5937031484257871
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 1.0}
    train score: 0.6061776061776061 validation score: 0.5919282511210762
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 1.0}
    train score: 0.6064703042008691 validation score: 0.5877061469265367
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 1.0}
    train score: 0.6074360212457751 validation score: 0.5877061469265367
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 1.0}
    train score: 0.6061776061776061 validation score: 0.5877061469265367
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 1.0}
    train score: 0.6061776061776061 validation score: 0.5877061469265367
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 1.0}
    train score: 0.6064703042008691 validation score: 0.5877061469265367
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 10.0}
    train score: 0.6098265895953757 validation score: 0.5898203592814371
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 10.0}
    train score: 0.6098265895953757 validation score: 0.5898203592814371
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 10.0}
    train score: 0.6095329802599905 validation score: 0.5898203592814371
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 10.0}
    train score: 0.6095329802599905 validation score: 0.5898203592814371
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 10.0}
    train score: 0.6081927710843373 validation score: 0.5898203592814371
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 10.0}
    train score: 0.6081927710843373 validation score: 0.5898203592814371
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 10.0}
    train score: 0.6081927710843373 validation score: 0.5898203592814371
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 100.0}
    train score: 0.6091566265060242 validation score: 0.5898203592814371
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 100.0}
    train score: 0.6091566265060242 validation score: 0.5898203592814371
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':

```



```

0.75, 'C': 10000.0}
    train score: 0.6091566265060242 validation score: 0.5898203592814371
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 10000.0}
    train score: 0.6091566265060242 validation score: 0.5898203592814371
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 10000.0}
    train score: 0.6091566265060242 validation score: 0.5898203592814371
[28]
LogisticRegression(l1_ratio=0.01, max_iter=100000, penalty='elasticnet',
                    random_state=132, solver='saga')

best model parameters: {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter':
100000, 'l1_ratio': 0.01, 'C': 1.0}
corresponding validation F1 score: 0.5937031484257871
test F1 score: 0.5982658959537572

randoms state 8
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 0.001}
    train score: 0.20109976433621368 validation score: 0.15130023640661938
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 0.001}
    train score: 0.0 validation score: 0.0

```

```

    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 0.01}
    train score: 0.5837615621788282 validation score: 0.5195618153364632
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 0.01}
    train score: 0.5767634854771784 validation score: 0.5087440381558028
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 0.01}
    train score: 0.5655391120507399 validation score: 0.5080385852090032
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 0.01}
    train score: 0.5428881650380022 validation score: 0.5008130081300813
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 0.01}
    train score: 0.5345113197128658 validation score: 0.4943089430894309
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 0.01}
    train score: 0.5277161862527716 validation score: 0.4868421052631579
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 0.01}
    train score: 0.5245720596355604 validation score: 0.48344370860927144
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 0.1}
    train score: 0.6128873585833744 validation score: 0.5663716814159293
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 0.1}
    train score: 0.6131889763779528 validation score: 0.5663716814159293
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 0.1}
    train score: 0.6141732283464566 validation score: 0.5650887573964498
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 0.1}
    train score: 0.6151574803149606 validation score: 0.5621301775147929
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 0.1}
    train score: 0.613714849531327 validation score: 0.5557206537890045

```

```

    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 0.1}
    train score: 0.6153846153846154 validation score: 0.5514157973174367
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 0.1}
    train score: 0.6164451009354996 validation score: 0.5500747384155457
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 1.0}
    train score: 0.6159844054580896 validation score: 0.5726744186046511
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 1.0}
    train score: 0.6159844054580896 validation score: 0.5726744186046511
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 1.0}
    train score: 0.6159844054580896 validation score: 0.5726744186046511
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 1.0}
    train score: 0.6159102000976086 validation score: 0.5718432510885341
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 1.0}
    train score: 0.615609756097561 validation score: 0.5735080058224162
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 1.0}
    train score: 0.6159844054580896 validation score: 0.5735080058224162
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 1.0}
    train score: 0.6156843643448612 validation score: 0.5755813953488371
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 10.0}
    train score: 0.6159844054580896 validation score: 0.5797101449275363
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 10.0}
    train score: 0.6159844054580896 validation score: 0.5797101449275363
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 10.0}
    train score: 0.6159844054580896 validation score: 0.5797101449275363
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 10.0}
    train score: 0.6162847391516334 validation score: 0.5797101449275363
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 10.0}
    train score: 0.6162847391516334 validation score: 0.5797101449275363
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 10.0}
    train score: 0.6162847391516334 validation score: 0.5797101449275363
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 10.0}
    train score: 0.6162847391516334 validation score: 0.5797101449275363

```

```

    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 100.0}
    train score: 0.6160583941605839 validation score: 0.5817655571635311
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 100.0}
    train score: 0.6160583941605839 validation score: 0.5817655571635311
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 100.0}
    train score: 0.6160583941605839 validation score: 0.5817655571635311
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 100.0}
    train score: 0.6160583941605839 validation score: 0.5817655571635311
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 100.0}
    train score: 0.6153846153846154 validation score: 0.5797101449275363
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 100.0}
    train score: 0.6153846153846154 validation score: 0.5797101449275363
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 100.0}
    train score: 0.6153846153846154 validation score: 0.5797101449275363
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 1000.0}
    train score: 0.6160583941605839 validation score: 0.5838150289017341
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 1000.0}
    train score: 0.6160583941605839 validation score: 0.5838150289017341
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 1000.0}
    train score: 0.6160583941605839 validation score: 0.5838150289017341
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 1000.0}
    train score: 0.6160583941605839 validation score: 0.5838150289017341
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 1000.0}
    train score: 0.6160583941605839 validation score: 0.5838150289017341
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 1000.0}
    train score: 0.6160583941605839 validation score: 0.5838150289017341
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 1000.0}
    train score: 0.6160583941605839 validation score: 0.5838150289017341
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 10000.0}
    train score: 0.6160583941605839 validation score: 0.5838150289017341
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 10000.0}
    train score: 0.6160583941605839 validation score: 0.5838150289017341

```

```

    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 10000.0}
    train score: 0.6160583941605839 validation score: 0.5838150289017341
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 10000.0}
    train score: 0.6160583941605839 validation score: 0.5838150289017341
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 10000.0}
    train score: 0.6160583941605839 validation score: 0.5838150289017341
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 10000.0}
    train score: 0.6160583941605839 validation score: 0.5838150289017341
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 10000.0}
    train score: 0.6160583941605839 validation score: 0.5838150289017341

```

[49]

```

LogisticRegression(C=1000.0, l1_ratio=0.01, max_iter=100000,
                    penalty='elasticnet', random_state=154, solver='saga')

```

```

best model parameters: {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter':
100000, 'l1_ratio': 0.01, 'C': 1000.0}
corresponding validation F1 score: 0.5838150289017341
test F1 score: 0.5861561119293078

```

randoms state 9

```

    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 0.001}

```

```

train score: 0.1774960380348653 validation score: 0.14457831325301204
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 0.001}
train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 0.001}
train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 0.001}
train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 0.001}
train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 0.001}
train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 0.001}
train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 0.01}
train score: 0.561494551115724 validation score: 0.5670731707317073
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 0.01}
train score: 0.5547981122181437 validation score: 0.5687789799072644
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 0.01}
train score: 0.5344735435595938 validation score: 0.5615141955835963
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 0.01}
train score: 0.5189248491497532 validation score: 0.5138211382113822
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 0.01}
train score: 0.5094339622641509 validation score: 0.4966666666666666
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 0.01}
train score: 0.501952035694367 validation score: 0.4916387959866221
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 0.01}
train score: 0.4991587212563096 validation score: 0.4916387959866221
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 0.1}
train score: 0.5933528836754642 validation score: 0.5985185185185184
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 0.1}
train score: 0.5928466438020579 validation score: 0.5985185185185184
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 0.1}

```



```

train score: 0.5914664051005394 validation score: 0.5976331360946746
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 0.1}
train score: 0.5887573964497042 validation score: 0.5949926362297496
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 0.1}
train score: 0.5901477832512316 validation score: 0.5911764705882353
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 0.1}
train score: 0.5897435897435896 validation score: 0.5899705014749264
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 0.1}
train score: 0.5879446640316206 validation score: 0.587887740029542
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 1.0}
train score: 0.5991274842462433 validation score: 0.6075581395348837
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 1.0}
train score: 0.5991274842462433 validation score: 0.6055312954876273
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 1.0}
train score: 0.5981580222976249 validation score: 0.6034985422740524
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 1.0}
train score: 0.5984481086323958 validation score: 0.6034985422740524
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 1.0}
train score: 0.5974781765276429 validation score: 0.6055312954876273
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 1.0}
train score: 0.5988372093023256 validation score: 0.6034985422740524
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 1.0}
train score: 0.5988372093023256 validation score: 0.6034985422740524
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 10.0}
train score: 0.6008708272859217 validation score: 0.6075581395348837
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 10.0}
train score: 0.5999032414126754 validation score: 0.6075581395348837
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 10.0}
train score: 0.6001936108422071 validation score: 0.6075581395348837
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 10.0}
train score: 0.6001936108422071 validation score: 0.6075581395348837
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 10.0}

```

```

train score: 0.6004842615012106 validation score: 0.6075581395348837
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 10.0}
train score: 0.6004842615012106 validation score: 0.6075581395348837
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 10.0}
train score: 0.6004842615012106 validation score: 0.6075581395348837
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 100.0}
train score: 0.5999032414126754 validation score: 0.6055312954876273
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 100.0}
train score: 0.5999032414126754 validation score: 0.6055312954876273
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 100.0}
train score: 0.5999032414126754 validation score: 0.6055312954876273
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 100.0}
train score: 0.6005802707930368 validation score: 0.6055312954876273
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 100.0}
train score: 0.6005802707930368 validation score: 0.6055312954876273
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 100.0}
train score: 0.6005802707930368 validation score: 0.6055312954876273
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 100.0}
train score: 0.6005802707930368 validation score: 0.6055312954876273
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 1000.0}
train score: 0.5999032414126754 validation score: 0.6055312954876273
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 1000.0}
train score: 0.5999032414126754 validation score: 0.6055312954876273
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 1000.0}
train score: 0.5999032414126754 validation score: 0.6055312954876273
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 1000.0}
train score: 0.5999032414126754 validation score: 0.6055312954876273
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 1000.0}
train score: 0.5999032414126754 validation score: 0.6055312954876273
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 1000.0}
train score: 0.5999032414126754 validation score: 0.6055312954876273
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 1000.0}

```

```

train score: 0.5999032414126754 validation score: 0.6055312954876273
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 10000.0}
train score: 0.5999032414126754 validation score: 0.6055312954876273
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 10000.0}
train score: 0.5999032414126754 validation score: 0.6055312954876273
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 10000.0}
train score: 0.5999032414126754 validation score: 0.6055312954876273
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 10000.0}
train score: 0.5999032414126754 validation score: 0.6055312954876273
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 10000.0}
train score: 0.5999032414126754 validation score: 0.6055312954876273
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 10000.0}
train score: 0.5999032414126754 validation score: 0.6055312954876273
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 10000.0}
train score: 0.5999032414126754 validation score: 0.6055312954876273
[28]
LogisticRegression(l1_ratio=0.01, max_iter=100000, penalty='elasticnet',
                    random_state=176, solver='saga')

```

```

best model parameters: {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter':
100000, 'l1_ratio': 0.01, 'C': 1.0}
corresponding validation F1 score: 0.6075581395348837
test F1 score: 0.5941176470588235

```

```

randoms state 10
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 0.0001}
train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 0.0001}
train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 0.0001}
train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 0.0001}
train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 0.0001}
train score: 0.0 validation score: 0.0
{'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':

```

```

0.9, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 0.0001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 0.001}
    train score: 0.18196202531645572 validation score: 0.1674641148325359
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 0.001}
    train score: 0.0 validation score: 0.0
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 0.01}
    train score: 0.5674273858921163 validation score: 0.5454545454545454
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 0.01}
    train score: 0.566839378238342 validation score: 0.5352112676056339
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 0.01}
    train score: 0.5577227200843436 validation score: 0.5161290322580646
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 0.01}
    train score: 0.53470715835141 validation score: 0.5057096247960848
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 0.01}
    train score: 0.5263157894736842 validation score: 0.5
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 0.01}
    train score: 0.519580805295091 validation score: 0.5
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 0.01}
    train score: 0.5176211453744494 validation score: 0.4983498349834984
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':

```

```

0.01, 'C': 0.1}
    train score: 0.604197169350903 validation score: 0.577259475218659
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 0.1}
    train score: 0.6013712047012733 validation score: 0.5789473684210527
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 0.1}
    train score: 0.6000978952520802 validation score: 0.5797950219619327
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 0.1}
    train score: 0.5998043052837574 validation score: 0.5747800586510263
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 0.1}
    train score: 0.5950738916256156 validation score: 0.5726872246696034
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 0.1}
    train score: 0.5952615992102666 validation score: 0.5731166912850812
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 0.1}
    train score: 0.5951557093425606 validation score: 0.5718518518518518
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 1.0}
    train score: 0.6092843326885881 validation score: 0.5820256776034237
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 1.0}
    train score: 0.6095791001451378 validation score: 0.5820256776034237
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 1.0}
    train score: 0.6095791001451378 validation score: 0.5820256776034237
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 1.0}
    train score: 0.6101694915254238 validation score: 0.5820256776034237
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 1.0}
    train score: 0.6092009685230024 validation score: 0.5820256776034237
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 1.0}
    train score: 0.6085271317829458 validation score: 0.5820256776034237
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 1.0}
    train score: 0.6085271317829458 validation score: 0.5820256776034237
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 10.0}
    train score: 0.6073500967117988 validation score: 0.582857142857143
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 10.0}
    train score: 0.6070565490575157 validation score: 0.582857142857143
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':

```



```

0.75, 'C': 1000.0}
    train score: 0.6070565490575157 validation score: 0.5820256776034237
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 1000.0}
    train score: 0.6070565490575157 validation score: 0.5820256776034237
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 1000.0}
    train score: 0.6070565490575157 validation score: 0.5820256776034237
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.01, 'C': 10000.0}
    train score: 0.6070565490575157 validation score: 0.5820256776034237
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.1, 'C': 10000.0}
    train score: 0.6070565490575157 validation score: 0.5820256776034237
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.25, 'C': 10000.0}
    train score: 0.6070565490575157 validation score: 0.5820256776034237
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.5, 'C': 10000.0}
    train score: 0.6070565490575157 validation score: 0.5820256776034237
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.75, 'C': 10000.0}
    train score: 0.6070565490575157 validation score: 0.5820256776034237
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.9, 'C': 10000.0}
    train score: 0.6070565490575157 validation score: 0.5820256776034237
    {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter': 100000, 'l1_ratio':
0.99, 'C': 10000.0}
    train score: 0.6070565490575157 validation score: 0.5820256776034237
[35]
LogisticRegression(C=10.0, l1_ratio=0.01, max_iter=100000, penalty='elasticnet',
                    random_state=198, solver='saga')

best model parameters: {'solver': 'saga', 'penalty': 'elasticnet', 'max_iter':
100000, 'l1_ratio': 0.01, 'C': 10.0}
corresponding validation F1 score: 0.582857142857143
test F1 score: 0.5988372093023255

```

```

[872]: # Saving Logistic Elastic Net Best Models at 10 Random States
filename = dirct + '/results/log_EN_best_models.sav'
joblib.dump(LogEN_models, filename)

```

```

[872]: ['/Users/liyuetian1/Documents/GitHub/DATA1030_MidtermProject/results/log_EN_best
_models.sav']

```

```

[342]: # SVC
from sklearn.svm import SVC

```

```

params = {
    'gamma' : [1e-2, 1e-1, 1e0, 1e1, 1e2, 'auto', 'scale'],
    'C': [0.01, 0.1, 0.5, 1, 5, 10, 20]
}

SVC_val_best_F1, SVC_test_F1, SVC_models = MLpipe_Stratify_f1(X, y,
    ↪preprocessor, SVC, params)

```

```

randoms state 1
[20]
SVC(C=0.5, random_state=0)

best model parameters: {'gamma': 'scale', 'C': 0.5}
corresponding validation F1 score: 0.5861027190332326
test F1 score: 0.5834633385335413

```

```

randoms state 2
[34]
SVC(C=5, random_state=22)

best model parameters: {'gamma': 'scale', 'C': 5}
corresponding validation F1 score: 0.554858934169279
test F1 score: 0.5578446909667196

```

```

randoms state 3
[34]
SVC(C=5, random_state=44)

best model parameters: {'gamma': 'scale', 'C': 5}
corresponding validation F1 score: 0.5892857142857142
test F1 score: 0.5795275590551182

```

```

randoms state 4
[35]
SVC(C=10, gamma=0.01, random_state=66)

best model parameters: {'gamma': 0.01, 'C': 10}
corresponding validation F1 score: 0.5710059171597633
test F1 score: 0.5628930817610063

```

```

randoms state 5
[29]
SVC(C=5, gamma=0.1, random_state=88)

best model parameters: {'gamma': 0.1, 'C': 5}

```


corresponding validation F1 score: 0.5910447761194031
test F1 score: 0.573082489146165

randoms state 6

[34]

SVC(C=5, random_state=110)

best model parameters: {'gamma': 'scale', 'C': 5}
corresponding validation F1 score: 0.5885885885885886
test F1 score: 0.5592705167173253

randoms state 7

[29]

SVC(C=5, gamma=0.1, random_state=132)

best model parameters: {'gamma': 0.1, 'C': 5}
corresponding validation F1 score: 0.5585585585585586
test F1 score: 0.5756676557863502

randoms state 8

[41]

SVC(C=10, random_state=154)

best model parameters: {'gamma': 'scale', 'C': 10}
corresponding validation F1 score: 0.5567010309278351
test F1 score: 0.5666666666666668

randoms state 9

[34]

SVC(C=5, random_state=176)

best model parameters: {'gamma': 'scale', 'C': 5}
corresponding validation F1 score: 0.5785609397944199
test F1 score: 0.5603715170278637

randoms state 10

[47]

SVC(C=20, gamma='auto', random_state=198)

best model parameters: {'gamma': 'auto', 'C': 20}
corresponding validation F1 score: 0.5641791044776119
test F1 score: 0.5722983257229831

[873]: *# Saving SVC Best Models at 10 Random States*

```
filename = direct + '/results/SVC_best_models.sav'  
joblib.dump(SVC_models, filename)
```

```
[873]: ['/Users/liyuetian1/Documents/GitHub/DATA1030_MidtermProject/results/SVC_best_models.sav']
```

```
[340]: # KNN

from sklearn.neighbors import KNeighborsClassifier

params = {
    'n_neighbors': [1, 2, 3, 5, 10, 30, 50, 100, 200],
    'weights': ['uniform', 'distance']
}

KNN_val_best_F1, KNN_test_F1, KNN_models = MLpipe_Stratify_f1(X, y,
    ↪preprocessor, KNeighborsClassifier, params)
```

randoms state 1

[17]

KNeighborsClassifier(n_neighbors=200, weights='distance')

best model parameters: {'weights': 'distance', 'n_neighbors': 200}

corresponding validation F1 score: 0.5888594164456235

test F1 score: 0.5897079276773296

randoms state 2

[10]

KNeighborsClassifier(n_neighbors=30)

best model parameters: {'weights': 'uniform', 'n_neighbors': 30}

corresponding validation F1 score: 0.5688888888888889

test F1 score: 0.5901162790697674

randoms state 3

[17]

KNeighborsClassifier(n_neighbors=200, weights='distance')

best model parameters: {'weights': 'distance', 'n_neighbors': 200}

corresponding validation F1 score: 0.5875862068965517

test F1 score: 0.6092124814264487

randoms state 4

[12]

KNeighborsClassifier(n_neighbors=50)

best model parameters: {'weights': 'uniform', 'n_neighbors': 50}

corresponding validation F1 score: 0.5969738651994497

test F1 score: 0.5760233918128655

```

randoms state 5
[12]
KNeighborsClassifier(n_neighbors=50)

best model parameters: {'weights': 'uniform', 'n_neighbors': 50}
corresponding validation F1 score: 0.593886462882096
test F1 score: 0.5932203389830508

randoms state 6
[17]
KNeighborsClassifier(n_neighbors=200, weights='distance')

best model parameters: {'weights': 'distance', 'n_neighbors': 200}
corresponding validation F1 score: 0.6
test F1 score: 0.5694249649368864

randoms state 7
[16]
KNeighborsClassifier(n_neighbors=200)

best model parameters: {'weights': 'uniform', 'n_neighbors': 200}
corresponding validation F1 score: 0.6008583690987125
test F1 score: 0.5573294629898403

randoms state 8
[14]
KNeighborsClassifier(n_neighbors=100)

best model parameters: {'weights': 'uniform', 'n_neighbors': 100}
corresponding validation F1 score: 0.5751072961373391
test F1 score: 0.5714285714285715

randoms state 9
[12]
KNeighborsClassifier(n_neighbors=50)

best model parameters: {'weights': 'uniform', 'n_neighbors': 50}
corresponding validation F1 score: 0.6067415730337079
test F1 score: 0.5965417867435159

randoms state 10
[17]
KNeighborsClassifier(n_neighbors=200, weights='distance')

best model parameters: {'weights': 'distance', 'n_neighbors': 200}
corresponding validation F1 score: 0.5726256983240224
test F1 score: 0.6102635228848821

```

```
[874]: # Saving KNN Best Models at 10 Random States
```

```
filename = direct + '/results/KNN_best_models.sav'  
joblib.dump(KNN_models, filename)
```

```
[874]: ['/Users/liyuetian1/Documents/GitHub/DATA1030_MidtermProject/results/KNN_best_models.sav']
```

```
[341]: # Random Forest
```

```
from sklearn.ensemble import RandomForestClassifier  
  
params = { 'max_features': [1, 3, 5, 10, 20, None],  
           'max_depth': [1, 3, 5, 7, 10, 15, 20, None]}  
          # 'min_samples_split': [ 2, 5, 10, 15, 25, 30] }  
  
RF_val_best_F1, RF_test_F1, RF_models = MLpipe_Stratify_f1(X, y, preprocessor,   
↳ RandomForestClassifier, params)
```

```
randoms state 1
```

```
[138]
```

```
RandomForestClassifier(max_depth=7, max_features=None, random_state=0)
```

```
best model parameters: {'min_samples_split': 2, 'max_features': None,  
'max_depth': 7}
```

```
corresponding validation F1 score: 0.596045197740113
```

```
test F1 score: 0.5964391691394659
```

```
randoms state 2
```

```
[192]
```

```
RandomForestClassifier(max_depth=15, max_features=5, random_state=22)
```

```
best model parameters: {'min_samples_split': 2, 'max_features': 5, 'max_depth':  
15}
```

```
corresponding validation F1 score: 0.5792507204610953
```

```
test F1 score: 0.5814307458143075
```

```
randoms state 3
```

```
[121]
```

```
RandomForestClassifier(max_depth=7, max_features=5, min_samples_split=5,  
                        random_state=44)
```

```
best model parameters: {'min_samples_split': 5, 'max_features': 5, 'max_depth':  
7}
```

```
corresponding validation F1 score: 0.5852187028657617
```

```
test F1 score: 0.577922077922078
```

```

randoms state 4
[132]
RandomForestClassifier(max_depth=7, max_features=20, random_state=66)

best model parameters: {'min_samples_split': 2, 'max_features': 20, 'max_depth':
7}
corresponding validation F1 score: 0.5891016200294551
test F1 score: 0.5740458015267177

randoms state 5
[130]
RandomForestClassifier(max_depth=7, max_features=10, min_samples_split=25,
                        random_state=88)

best model parameters: {'min_samples_split': 25, 'max_features': 10,
'max_depth': 7}
corresponding validation F1 score: 0.5899705014749264
test F1 score: 0.5944363103953147

randoms state 6
[209]
RandomForestClassifier(max_depth=15, max_features=20, min_samples_split=30,
                        random_state=110)

best model parameters: {'min_samples_split': 30, 'max_features': 20,
'max_depth': 15}
corresponding validation F1 score: 0.6156069364161849
test F1 score: 0.5692995529061102

randoms state 7
[142]
RandomForestClassifier(max_depth=7, max_features=None, min_samples_split=25,
                        random_state=132)

best model parameters: {'min_samples_split': 25, 'max_features': None,
'max_depth': 7}
corresponding validation F1 score: 0.5786350148367952
test F1 score: 0.5718518518518518

randoms state 8
[158]
RandomForestClassifier(max_depth=10, max_features=5, min_samples_split=10,
                        random_state=154)

best model parameters: {'min_samples_split': 10, 'max_features': 5, 'max_depth':
10}
corresponding validation F1 score: 0.562406015037594

```

```
test F1 score: 0.576271186440678
```

```
randoms state 9
```

```
[127]
```

```
RandomForestClassifier(max_depth=7, max_features=10, min_samples_split=5,  
                        random_state=176)
```

```
best model parameters: {'min_samples_split': 5, 'max_features': 10, 'max_depth':  
7}
```

```
corresponding validation F1 score: 0.5889387144992526
```

```
test F1 score: 0.5875190258751902
```

```
randoms state 10
```

```
[170]
```

```
RandomForestClassifier(max_depth=10, max_features=20, min_samples_split=10,  
                        random_state=198)
```

```
best model parameters: {'min_samples_split': 10, 'max_features': 20,  
'max_depth': 10}
```

```
corresponding validation F1 score: 0.5681159420289854
```

```
test F1 score: 0.5892857142857142
```

```
[875]: # Saving RF Best Models at 10 Random States
```

```
filename = direct + '/results/RF_best_models.sav'  
joblib.dump(RF_models, filename)
```

```
[875]: ['/Users/liyuetian1/Documents/GitHub/DATA1030_MidtermProject/results/RF_best_mod  
els.sav']
```

```
[359]: # Xgboost
```

```
from sklearn.model_selection import ParameterGrid  
import xgboost
```

```
param_grid = {  
    "learning_rate": [0.05, 0.1, 0.2, 0.3],  
    "n_estimators": [1000],  
    "seed": [0],  
    "gamma": [0, 0.1, 0.2, 0.3, 0.4],  
    "colsample_bytree": [0.3, 0.4, 0.5, 0.7],  
    "subsample": [0.4, 0.5, 0.65, 0.75, 1],  
    "min_child_weight": [1, 3, 5, 7],  
    "eval_metric": ['logloss']  
}
```

```
nr_states = 10
```

```

XGB_test_scores = np.zeros(nr_states)
XGB_val_best_scores = np.zeros(nr_states)
XGB_final_models = []
XGB_feature_importances = np.zeros(nr_states)

for i in range(nr_states):

    models = []

    print('\nrandoms state '+str(i+1))

    X_other, X_test, y_other, y_test = train_test_split(X,y,test_size = 0.
→2, stratify = y, random_state=22*i)
    X_train, X_val, y_train, y_val = train_test_split(X_other,y_other,test_size_
→= 0.25, stratify = y_other, random_state=22*i)

    ## Preprocess
    X_train_prep = preprocessor.fit_transform(X_train)
    X_val_prep = preprocessor.transform(X_val)
    X_test_prep = preprocessor.transform(X_test)

    XGB_feature_names = preprocessor.transformers_[0][-1] + \
        list(preprocessor.named_transformers_['cat'][0].
→get_feature_names(cat_fts)) + \
        preprocessor.transformers_[2][-1]

    train_score = np.zeros(len(ParameterGrid(param_grid)))
    val_score = np.zeros(len(ParameterGrid(param_grid)))

    ## Loop through parameters
    for p in range(len(ParameterGrid(param_grid))):
        params = ParameterGrid(param_grid)[p]
        XGB = xgboost.XGBClassifier(use_label_encoder =False, random_state =
→22*i)
        XGB.set_params(**params)
        XGB.fit(X_train_prep,y_train,early_stopping_rounds=50,
→eval_set= [(X_val_prep, y_val)], verbose=False)

        models.append(XGB) # save it
        y_val_pred = XGB.predict(X_val_prep)

        val_score[p] = f1_score(y_val, y_val_pred)

    XGB_val_best_scores = np.max(val_score)

    print('\nbest model parameters:',ParameterGrid(param_grid)[np.
→argmax(val_score)])

```

```

print('corresponding validation score:', np.max(val_score))

XGB_final_models.append(models[np.argmax(val_score)])

y_test_pred = XGB_final_models[-1].predict(X_test_prep)

# calculate and save the test score
XBG_test_scores[i] = f1_score(y_test, y_test_pred)

print('test f1 score:', XBG_test_scores[i])

```

randoms state 1

```

best model parameters: {'subsample': 0.4, 'seed': 0, 'n_estimators': 1000,
'min_child_weight': 1, 'learning_rate': 0.2, 'gamma': 0.2, 'eval_metric':
'logloss', 'colsample_bytree': 0.5}
corresponding validation score: 0.6031294452347084
test f1 score: 0.5882352941176471

```

randoms state 2

```

best model parameters: {'subsample': 0.5, 'seed': 0, 'n_estimators': 1000,
'min_child_weight': 7, 'learning_rate': 0.3, 'gamma': 0.3, 'eval_metric':
'logloss', 'colsample_bytree': 0.3}
corresponding validation score: 0.5937031484257871
test f1 score: 0.5727554179566563

```

randoms state 3

```

best model parameters: {'subsample': 0.4, 'seed': 0, 'n_estimators': 1000,
'min_child_weight': 5, 'learning_rate': 0.2, 'gamma': 0.3, 'eval_metric':
'logloss', 'colsample_bytree': 0.3}
corresponding validation score: 0.6023391812865497
test f1 score: 0.5750000000000001

```

randoms state 4

```

best model parameters: {'subsample': 0.65, 'seed': 0, 'n_estimators': 1000,
'min_child_weight': 3, 'learning_rate': 0.3, 'gamma': 0.1, 'eval_metric':
'logloss', 'colsample_bytree': 0.7}
corresponding validation score: 0.6054519368723099
test f1 score: 0.5630498533724341

```

randoms state 5

best model parameters: {'subsample': 1, 'seed': 0, 'n_estimators': 1000,
'min_child_weight': 5, 'learning_rate': 0.3, 'gamma': 0.1, 'eval_metric':
'logloss', 'colsample_bytree': 0.3}
corresponding validation score: 0.5994152046783625
test f1 score: 0.6051873198847263

randoms state 6

best model parameters: {'subsample': 0.65, 'seed': 0, 'n_estimators': 1000,
'min_child_weight': 3, 'learning_rate': 0.1, 'gamma': 0.1, 'eval_metric':
'logloss', 'colsample_bytree': 0.7}
corresponding validation score: 0.6251808972503619
test f1 score: 0.5837037037037036

randoms state 7

best model parameters: {'subsample': 0.4, 'seed': 0, 'n_estimators': 1000,
'min_child_weight': 5, 'learning_rate': 0.3, 'gamma': 0.4, 'eval_metric':
'logloss', 'colsample_bytree': 0.5}
corresponding validation score: 0.6002886002886003
test f1 score: 0.6025824964131994

randoms state 8

best model parameters: {'subsample': 0.4, 'seed': 0, 'n_estimators': 1000,
'min_child_weight': 5, 'learning_rate': 0.2, 'gamma': 0.3, 'eval_metric':
'logloss', 'colsample_bytree': 0.7}
corresponding validation score: 0.5852941176470589
test f1 score: 0.5612403100775193

randoms state 9

best model parameters: {'subsample': 0.75, 'seed': 0, 'n_estimators': 1000,
'min_child_weight': 1, 'learning_rate': 0.3, 'gamma': 0.3, 'eval_metric':
'logloss', 'colsample_bytree': 0.5}
corresponding validation score: 0.6075581395348837
test f1 score: 0.5843373493975903

randoms state 10

best model parameters: {'subsample': 0.5, 'seed': 0, 'n_estimators': 1000,
'min_child_weight': 1, 'learning_rate': 0.3, 'gamma': 0.4, 'eval_metric':
'logloss', 'colsample_bytree': 0.3}
corresponding validation score: 0.5906432748538012
test f1 score: 0.5825825825825826

```
[877]: # Saving RF Best Models at 10 Random States
```


```
filename = direct + '/results/XGBoost_best_models.sav'  
joblib.dump(XGB_final_models, filename)
```

```
[877]: ['/Users/liyuetian1/Documents/GitHub/DATA1030_MidtermProject/results/XGBoost_best_models.sav']
```

3 3. Results

```
[884]: # Models F1 Scores Summary
```

```
rand_states = np.arange(1,11,1)  
  
column_names = ['L1', 'L2', 'EN', 'SVC', 'KNN', 'RF', 'XGBoost']  
  
models_F1 = pd.DataFrame([Logl1_test_F1, Logl2_test_F1, LogEN_test_F1,   
    ↳ SVC_test_F1, KNN_test_F1, RF_test_F1, XBG_test_scores], index=column_names,   
    ↳ columns = rand_states)  
  
models_F1 = models_F1.T  
  
models_F1  
  
mean_F1 = models_F1.mean(axis=0)  
std_F1 = models_F1.std(axis=0)  
  
plt.figure()  
plt.rcParams.update({'font.size': 13})  
plt.plot(column_names, mean_F1, 'o', color='k')  
plt.errorbar(column_names, mean_F1, yerr= std_F1, fmt='o', capsize= 0.5,   
    ↳ capthick= 0.2, ecolor='k')  
  
plt.ylabel('F1 Scores')  
plt.xlabel('ML Models')  
plt.title('Models F1 Scores')  
plt.xticks(rotation=20)  
plt.savefig(direct + '/figures/Models_F1_Scores.png', dpi=300)  
plt.show()
```



```
[364]: mean_F1 = models_F1.mean(axis=0)
       print(mean_F1)
```

```
[364]: Lasso      0.596412
       Ridge      0.596825
       EN         0.596300
       SVC        0.569109
       KNN        0.586327
       RF         0.581850
       XGBoost    0.581867
       dtype: float64
```

```
[879]: print(std_F1)
```

```
Lasso      0.008797
Ridge      0.008388
EN         0.008544
SVC        0.008995
KNN        0.017413
RF         0.009567
XGBoost    0.014656
dtype: float64
```

Given that Logistic L2 Regularization Model achieved the highest averaged F1 score with the second lowest standard deviations, it was chosen as the final model.

4 4. Model Interpretation

```
[474]: # Pipeline for Logistic Regression L2 Regularization Only

[36]: from sklearn.metrics import f1_score
from sklearn.metrics import accuracy_score
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.model_selection import ParameterGrid

def Logl2_Stratify_f1(X, y, preprocessor, ML_algo, param_grid, verbose = 1):
    '''
        This function intends to focus on analyzing and interpreting the final
        ↪model
        Logistic Regression L2 Regularization;
        - Same random states, parameters as previous one
        - Additional outputs for feature importance calculation
    '''

    nr_states = 10
    test_scores = np.zeros(nr_states)
    val_best_scores = np.zeros(nr_states)
    final_models = []
    feature_importances = np.zeros(nr_states)
    X_test_all = []
    Y_test_all = []

    for i in range(nr_states):

        print('\nrandoms state '+str(i+1))

        X_other, X_test, y_other, y_test = train_test_split(X,y,test_size = 0.
        ↪2,stratify = y, random_state=22*i)
        X_train, X_val, y_train, y_val =
        ↪train_test_split(X_other,y_other,test_size = 0.25, stratify = y_other,
        ↪random_state=22*i)

        train_score = np.zeros(len(ParameterGrid(param_grid)))
        val_score = np.zeros(len(ParameterGrid(param_grid)))

        X_train_prep = preprocessor.fit_transform(X_train)
```

```

feature_names = preprocessor.transformers_[0][-1] + \
    list(preprocessor.named_transformers_['cat'][0].
→get_feature_names(cat_ftrs)) + \
    preprocessor.transformers_[2][-1]

X_val_prep = preprocessor.transform(X_val)
X_test_prep = preprocessor.transform(X_test)

if verbose ==2 :
    final_scaler = StandardScaler()
    X_train_prep = final_scaler.fit_transform(X_train_prep)
    X_val_prep = final_scaler.transform(X_val_prep)
    X_test_prep = final_scaler.transform(X_test_prep)
    print('Mean Standardized All Features')

X_test_all.append(X_test_prep)
Y_test_all.append(y_test)

models = []

for p in range(len(ParameterGrid(param_grid))):
    params = ParameterGrid(param_grid)[p]
    try:
        ML = ML_algo(random_state = 22*i)
    except:
        ML = ML_algo()

    ML.set_params(**params)
    ML.fit(X_train_prep,y_train)

    train_score[p] = f1_score(y_train, ML.predict(X_train_prep))

    models.append(ML)
    y_CV_pred = ML.predict(X_val_prep)
    val_score[p]= f1_score(y_val, y_CV_pred)

    print([np.argmax(val_score)])
    print(models[np.argmax(val_score)])
    val_best_scores[i] = np.max(val_score)
    print('\nbest model parameters:',ParameterGrid(param_grid)[np.
→argmax(val_score)])
    print('corresponding validation F1 score:',np.max(val_score))

    final_models.append(models[np.argmax(val_score)])

```

```

y_test_pred = final_models[-1].predict(X_test_prep)

test_scores[i] = f1_score(y_test, y_test_pred)

print('test F1 score:',test_scores[i])

return val_best_scores, test_scores, final_models, X_test_all, Y_test_all,
↪feature_names

```

```

[38]: # Logistic Regression l2
      ## Verbose = 1, no final mean standardization

      from sklearn.linear_model import LogisticRegression

      params = { 'penalty' : ['l2'],
                  'C' : [1e-4, 1e-3, 1e-2, 1e-1, 1e0, 1e1, 1e2, 1e3, 1e4],
                  'max_iter': [10000],
                  'solver': ['saga'] }

      Logl2_val_best_F1, Logl2_test_F1_v1, Logl2_models_v1, all_X_test_v1,
      ↪all_Y_test_v1, feature_names = Logl2_Stratify_f1(X, y, processor,
      ↪LogisticRegression, params, 1)

```

randoms state 1

[5]

```
LogisticRegression(C=10.0, max_iter=10000, random_state=0, solver='saga')
```

```
best model parameters: {'solver': 'saga', 'penalty': 'l2', 'max_iter': 10000,
'C': 10.0}
```

```
corresponding validation F1 score: 0.5957446808510639
```

```
test F1 score: 0.6169590643274854
```

randoms state 2

[5]

```
LogisticRegression(C=10.0, max_iter=10000, random_state=22, solver='saga')
```

```
best model parameters: {'solver': 'saga', 'penalty': 'l2', 'max_iter': 10000,
'C': 10.0}
```

```
corresponding validation F1 score: 0.5819793205317577
```

```
test F1 score: 0.5889387144992526
```

randoms state 3

[3]

```
LogisticRegression(C=0.1, max_iter=10000, random_state=44, solver='saga')
```

best model parameters: {'solver': 'saga', 'penalty': 'l2', 'max_iter': 10000, 'C': 0.1}
corresponding validation F1 score: 0.5936599423631125
test F1 score: 0.5969230769230769

randoms state 4

[6]

LogisticRegression(C=100.0, max_iter=10000, random_state=66, solver='saga')

best model parameters: {'solver': 'saga', 'penalty': 'l2', 'max_iter': 10000, 'C': 100.0}
corresponding validation F1 score: 0.592489568845619
test F1 score: 0.5915080527086385

randoms state 5

[6]

LogisticRegression(C=100.0, max_iter=10000, random_state=88, solver='saga')

best model parameters: {'solver': 'saga', 'penalty': 'l2', 'max_iter': 10000, 'C': 100.0}
corresponding validation F1 score: 0.5852941176470589
test F1 score: 0.5965417867435159

randoms state 6

[5]

LogisticRegression(C=10.0, max_iter=10000, random_state=110, solver='saga')

best model parameters: {'solver': 'saga', 'penalty': 'l2', 'max_iter': 10000, 'C': 10.0}
corresponding validation F1 score: 0.6037735849056604
test F1 score: 0.6

randoms state 7

[4]

LogisticRegression(max_iter=10000, random_state=132, solver='saga')

best model parameters: {'solver': 'saga', 'penalty': 'l2', 'max_iter': 10000, 'C': 1.0}
corresponding validation F1 score: 0.5937031484257871
test F1 score: 0.5982658959537572

randoms state 8

[7]

LogisticRegression(C=1000.0, max_iter=10000, random_state=154, solver='saga')

best model parameters: {'solver': 'saga', 'penalty': 'l2', 'max_iter': 10000, 'C': 1000.0}
corresponding validation F1 score: 0.5838150289017341

```
test F1 score: 0.5861561119293078
```

```
randoms state 9
```

```
[4]
```

```
LogisticRegression(max_iter=10000, random_state=176, solver='saga')
```

```
best model parameters: {'solver': 'saga', 'penalty': 'l2', 'max_iter': 10000, 'C': 1.0}
```

```
corresponding validation F1 score: 0.6075581395348837
```

```
test F1 score: 0.5941176470588235
```

```
randoms state 10
```

```
[5]
```

```
LogisticRegression(C=10.0, max_iter=10000, random_state=198, solver='saga')
```

```
best model parameters: {'solver': 'saga', 'penalty': 'l2', 'max_iter': 10000, 'C': 10.0}
```

```
corresponding validation F1 score: 0.582857142857143
```

```
test F1 score: 0.5988372093023255
```

4.1 4.1 Global Feature Importance

4.1.1 4.1.1 Permutation Feature Importance

```
[39]: # Choose the Last Model from 10 Random States, C = 10.0
```

```
model = Logl2_models_v1[-1]
```

```
X_test_df = pd.DataFrame(data=all_X_test_v1[-1], columns=feature_names)
```

```
Y_test = all_Y_test_v1[-1]
```

```
np.random.seed(42)
```

```
nr_runs = 10
```

```
scores = np.zeros([len(feature_names),nr_runs])
```

```
for i in range(len(feature_names)):
```

```
    print('shuffling '+str(feature_names[i]))
```

```
    f1_scores = []
```

```
    for j in range(nr_runs):
```

```
        X_test_shuffled = X_test_df.copy()
```

```
        X_test_shuffled[feature_names[i]] = np.random.
```

```
        ↪ permutation(X_test_df[feature_names[i]].values)
```

```
        perm_Y_test_pred = model.predict(X_test_shuffled)
```

```
        f1_scores.append(f1_score(Y_test,perm_Y_test_pred))
```



```

    print('    shuffled test score:', np.around(np.mean(f1_scores), 3), '+/-' , np.
    ↪ around(np.std(f1_scores), 3))
    scores[i] = f1_scores

sorted_indcs = np.argsort(np.mean(scores, axis=1))[:, :-1]

plt.rcParams.update({'font.size': 14})
plt.figure(figsize=(10, 9))

label_ft = [feature_names[i] for i in sorted_indcs]

plt.boxplot(scores[sorted_indcs].T, labels= label_ft , vert=False)
plt.axvline(LogL2_test_F1_v1[-1], label='test F1 score')
plt.title("Permutation Importances (test set at random state 198)")
plt.xlabel('F1 score with perturbed feature')
plt.legend()
plt.tight_layout()
plt.savefig('figures/PermutationImportances_LogL2.png', dpi=300)
plt.show()

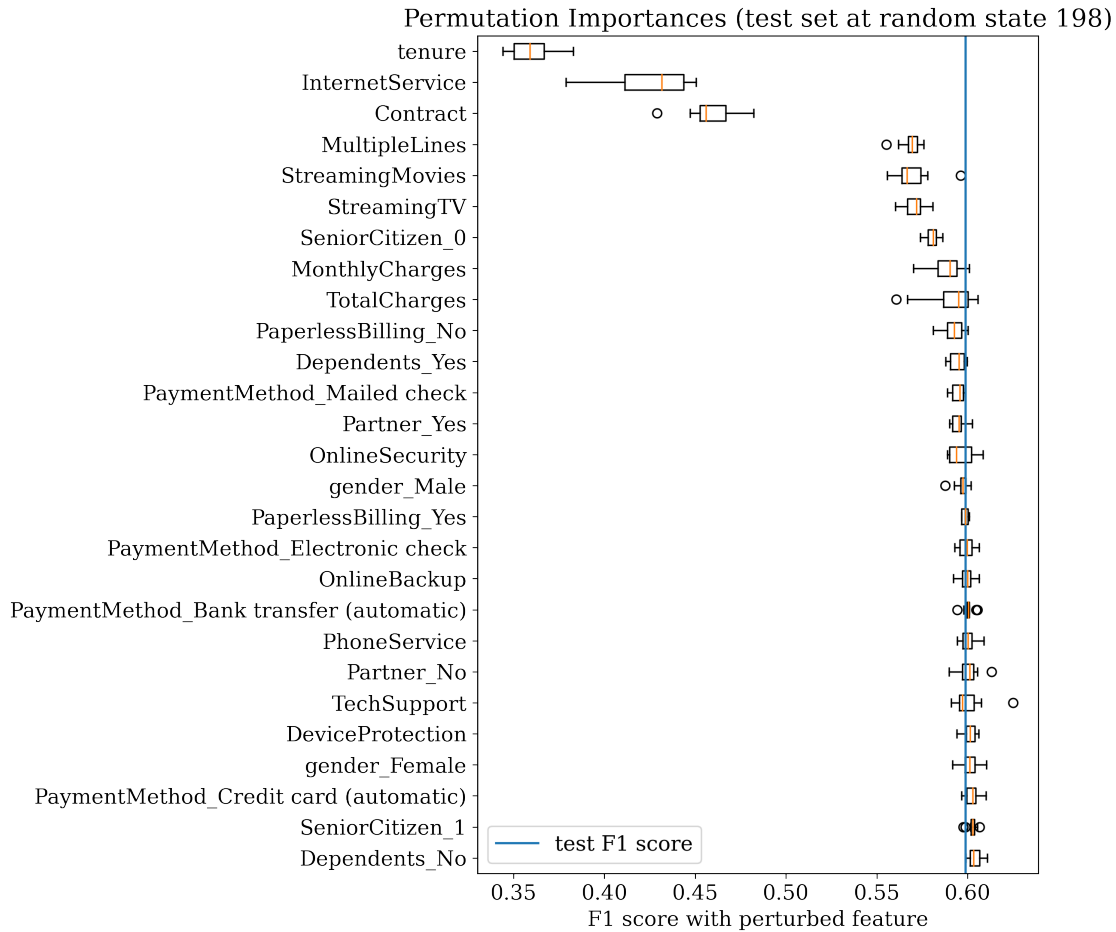
```

```

shuffling tenure
    shuffled test score: 0.359 +/- 0.012
shuffling MonthlyCharges
    shuffled test score: 0.588 +/- 0.008
shuffling TotalCharges
    shuffled test score: 0.59 +/- 0.014
shuffling gender_Female
    shuffled test score: 0.601 +/- 0.005
shuffling gender_Male
    shuffled test score: 0.597 +/- 0.004
shuffling SeniorCitizen_0
    shuffled test score: 0.581 +/- 0.004
shuffling SeniorCitizen_1
    shuffled test score: 0.603 +/- 0.003
shuffling Partner_No
    shuffled test score: 0.601 +/- 0.006
shuffling Partner_Yes
    shuffled test score: 0.595 +/- 0.004
shuffling Dependents_No
    shuffled test score: 0.604 +/- 0.004
shuffling Dependents_Yes
    shuffled test score: 0.594 +/- 0.004
shuffling PaperlessBilling_No
    shuffled test score: 0.592 +/- 0.005
shuffling PaperlessBilling_Yes
    shuffled test score: 0.599 +/- 0.002
shuffling PaymentMethod_Bank transfer (automatic)

```

shuffled test score: 0.6 +/- 0.003
shuffling PaymentMethod_Credit card (automatic)
shuffled test score: 0.602 +/- 0.004
shuffling PaymentMethod_Electronic check
shuffled test score: 0.599 +/- 0.004
shuffling PaymentMethod_Mailed check
shuffled test score: 0.595 +/- 0.003
shuffling PhoneService
shuffled test score: 0.601 +/- 0.004
shuffling MultipleLines
shuffled test score: 0.569 +/- 0.006
shuffling InternetService
shuffled test score: 0.425 +/- 0.023
shuffling OnlineSecurity
shuffled test score: 0.597 +/- 0.007
shuffling OnlineBackup
shuffled test score: 0.6 +/- 0.004
shuffling DeviceProtection
shuffled test score: 0.601 +/- 0.004
shuffling TechSupport
shuffled test score: 0.601 +/- 0.009
shuffling StreamingTV
shuffled test score: 0.571 +/- 0.007
shuffling StreamingMovies
shuffled test score: 0.57 +/- 0.011
shuffling Contract
shuffled test score: 0.458 +/- 0.014



4.1.2 Coefficients

```
[40]: # Verbose = 2, mean standardize all features to calculate scaled coefficients
```

```
params = { 'penalty' : ['l2'],
           'C' : [1e-4, 1e-3, 1e-2, 1e-1, 1e0, 1e1, 1e2, 1e3, 1e4],
           'max_iter': [10000],
           'solver': ['saga'] }

Logl2_val_best_F1, Logl2_test_F1, Logl2_models_v2, all_X_test_v2,
↳all_Y_test_v2, feature_names = Logl2_Stratify_f1(X, y, preprocessor,
↳LogisticRegression, params, 2)
```

randoms state 1

Mean Standardized All Features

[4]

LogisticRegression(max_iter=10000, random_state=0, solver='saga')

best model parameters: {'solver': 'saga', 'penalty': 'l2', 'max_iter': 10000, 'C': 1.0}
corresponding validation F1 score: 0.5968882602545968
test F1 score: 0.6218978102189782

randoms state 2

Mean Standardized All Features

[5]

LogisticRegression(C=10.0, max_iter=10000, random_state=22, solver='saga')

best model parameters: {'solver': 'saga', 'penalty': 'l2', 'max_iter': 10000, 'C': 10.0}
corresponding validation F1 score: 0.5819793205317577
test F1 score: 0.5880597014925374

randoms state 3

Mean Standardized All Features

[5]

LogisticRegression(C=10.0, max_iter=10000, random_state=44, solver='saga')

best model parameters: {'solver': 'saga', 'penalty': 'l2', 'max_iter': 10000, 'C': 10.0}
corresponding validation F1 score: 0.5894134477825466
test F1 score: 0.6079027355623101

randoms state 4

Mean Standardized All Features

[3]

LogisticRegression(C=0.1, max_iter=10000, random_state=66, solver='saga')

best model parameters: {'solver': 'saga', 'penalty': 'l2', 'max_iter': 10000, 'C': 0.1}
corresponding validation F1 score: 0.5938375350140056
test F1 score: 0.5920471281296025

randoms state 5

Mean Standardized All Features

[6]

LogisticRegression(C=100.0, max_iter=10000, random_state=88, solver='saga')

best model parameters: {'solver': 'saga', 'penalty': 'l2', 'max_iter': 10000, 'C': 100.0}
corresponding validation F1 score: 0.5852941176470589
test F1 score: 0.5965417867435159

randoms state 6

Mean Standardized All Features

```

[5]
LogisticRegression(C=10.0, max_iter=10000, random_state=110, solver='saga')

best model parameters: {'solver': 'saga', 'penalty': 'l2', 'max_iter': 10000,
'C': 10.0}
corresponding validation F1 score: 0.6037735849056604
test F1 score: 0.6

randoms state 7
Mean Standardized All Features
[4]
LogisticRegression(max_iter=10000, random_state=132, solver='saga')

best model parameters: {'solver': 'saga', 'penalty': 'l2', 'max_iter': 10000,
'C': 1.0}
corresponding validation F1 score: 0.5910447761194031
test F1 score: 0.5991316931982633

randoms state 8
Mean Standardized All Features
[6]
LogisticRegression(C=100.0, max_iter=10000, random_state=154, solver='saga')

best model parameters: {'solver': 'saga', 'penalty': 'l2', 'max_iter': 10000,
'C': 100.0}
corresponding validation F1 score: 0.5838150289017341
test F1 score: 0.5861561119293078

randoms state 9
Mean Standardized All Features
[4]
LogisticRegression(max_iter=10000, random_state=176, solver='saga')

best model parameters: {'solver': 'saga', 'penalty': 'l2', 'max_iter': 10000,
'C': 1.0}
corresponding validation F1 score: 0.6055312954876273
test F1 score: 0.5941176470588235

randoms state 10
Mean Standardized All Features
[4]
LogisticRegression(max_iter=10000, random_state=198, solver='saga')

best model parameters: {'solver': 'saga', 'penalty': 'l2', 'max_iter': 10000,
'C': 1.0}
corresponding validation F1 score: 0.5840455840455839
test F1 score: 0.6040462427745665

```

```
[41]: df = pd.DataFrame(columns=feature_names)
```

```
for i in range(len(Log12_models_v2)):
    coeffs = Log12_models_v2[i].coef_
    coeffs = coeffs.flatten()
    df.loc[len(df)] = coeffs

df
```

```
[41]:      tenure  MonthlyCharges  TotalCharges  gender_Female  gender_Male  \
0 -1.455692      -0.485221      0.717505      0.006505      -0.006505
1 -1.413223      -0.448889      0.653662      0.016476      -0.016476
2 -1.287476      -0.196579      0.572477      0.007304      -0.007304
3 -1.109718       0.028232      0.300720      0.005105      -0.005105
4 -1.400814      -0.508331      0.628667     -0.001046       0.001046
5 -1.722746      -0.592581      0.993347     -0.016378       0.016378
6 -1.446682      -0.483506      0.734251      0.010144     -0.010144
7 -1.624607      -0.151269      0.813818      0.001865     -0.001865
8 -1.352395      -0.481333      0.577631      0.015167     -0.015167
9 -1.410448      -0.540946      0.654020     -0.012774       0.012774

      SeniorCitizen_0  SeniorCitizen_1  Partner_No  Partner_Yes  Dependents_No  \
0      -0.038321      0.038321      0.010419     -0.010419      0.014601
1      -0.037852      0.037852      0.005839     -0.005839      0.028027
2      -0.025366      0.025366     -0.013296      0.013296      0.052123
3      -0.055956      0.055956     -0.023226      0.023226      0.065793
4      -0.033812      0.033812      0.002221     -0.002221      0.037855
5      -0.044401      0.044401     -0.004297      0.004297      0.014468
6      -0.048017      0.048017      0.028825     -0.028825      0.002484
7      -0.035895      0.035895     -0.003849      0.003849      0.037616
8      -0.041493      0.041493      0.003832     -0.003832      0.051598
9      -0.023258      0.023258      0.016375     -0.016375      0.044045

      ...  PhoneService  MultipleLines  InternetService  OnlineSecurity  \
0  ...      -0.119622      0.139245      0.832630      -0.184672
1  ...      -0.179400      0.245220      0.866304      -0.248488
2  ...      -0.304697      0.213082      0.729479      -0.260861
3  ...      -0.333867      0.169113      0.744774      -0.311952
4  ...      -0.208066      0.264237      0.952716      -0.209141
5  ...      -0.180441      0.195462      0.969734      -0.239209
6  ...      -0.154598      0.182867      0.937884      -0.278372
7  ...      -0.237823      0.140958      0.730981      -0.250504
8  ...      -0.187144      0.223166      0.864561      -0.164183
9  ...      -0.249395      0.263084      0.978672      -0.189443

      OnlineBackup  DeviceProtection  TechSupport  StreamingTV  StreamingMovies  \
0      -0.132269      0.074815     -0.263076      0.296624      0.254979
```

1	-0.104634	0.056365	-0.220938	0.281135	0.273190
2	-0.057228	0.043276	-0.277963	0.226318	0.245375
3	-0.082721	-0.044520	-0.135700	0.230002	0.131058
4	-0.048219	0.019298	-0.222804	0.276559	0.214199
5	-0.039704	-0.030931	-0.208715	0.261747	0.260644
6	-0.069856	-0.049884	-0.193395	0.301565	0.281122
7	-0.109290	-0.038859	-0.143211	0.222316	0.220335
8	-0.103697	0.078021	-0.273035	0.300060	0.253790
9	-0.043081	-0.044265	-0.197872	0.268159	0.271240

```

Contract
0 -0.605205
1 -0.602587
2 -0.576734
3 -0.552380
4 -0.585439
5 -0.594631
6 -0.568038
7 -0.539228
8 -0.547004
9 -0.617027

```

[10 rows x 27 columns]


```
[42]: cof_mean = df.mean(axis = 0)
      cof_std = df.std(axis = 0)
```

```
[43]: sorted_indcs = np.argsort(np.abs(cof_mean))

FN_coef = [feature_names[i] for i in sorted_indcs[-27:]]
FN_coef

plt.figure(figsize=(10,10))

plt.rcParams.update({'font.size': 14})
plt.barh(np.arange(27),cof_mean[sorted_indcs[-27:]], xerr = cof_std,
↪align='center', alpha=1.0, ecolor='black', capsize=5)
plt.yticks(np.arange(27),FN_coef)
plt.xlabel('coefficients')
plt.title('all scaled feature coefficients')
plt.tight_layout()
plt.savefig(direct + '/figures/LR_coefs_scaled.png',dpi=300)
plt.show()
```



4.1.3 Using SHAP to calculate global feature importance

```
[44]: import shap
      shap.initjs()
```

<IPython.core.display.HTML object>

```
[52]: # Use the 10th model at random state 22*9 to calculate shap values

y_test = all_Y_test_v1[-1]
x_test = pd.DataFrame(all_X_test_v1[-1], columns = feature_names)
model = Log12_models_v1[-1]
```




```


# Calculate Shap values for linear model
masker = shap.maskers.Independent(data=x_test, max_samples=1000)
explainer = shap.Explainer(model, masker=masker, feature_names= feature_names,
    ↪algorithm="linear")
shap_values = explainer.shap_values(x_test)

# Global Feature Importance based on Shap Values
shap.summary_plot(shap_values, x_test, plot_type="bar", feature_names=
    ↪feature_names, show=False)
plt.rcParams.update({'font.size': 10})
plt.savefig(direct + '/figures/Shap_Global_Feature_Importance.
    ↪png',bbox_inches='tight', dpi=250)

```



```
[51]: shap.summary_plot(shap_values, x_test)
```



This figure combines the feature importance and its coefficient values. Features are ordered based on their importances. For example, higher total charges make customers more likely to leave.


```
[61]: # Individual Feature: tenure

X_test_transformed = pd.DataFrame(x_test, columns = feature_names)
x_plot = X_test_transformed['tenure']

indx = feature_names.index('tenure')
indx

y_plot = shap_values[:,indx]
y_plot

matplotlib.rcParams.update({'font.size': 12})
plt.scatter(X_test_transformed['tenure'], shap_values[:,indx])
plt.ylabel('shap value')
plt.xlabel('tenure')
plt.show()
```



4.1.4 4.2 Local Feature Importance

```
[26]: # Taking the last model and test data set to look into local feature importance
      ↪ using Shap
```

```
[54]: y_pred = model.predict(x_test)
```

```
[915]: ## Looking at Costumer No.7

print(x_test.iloc[7])
print(shap_values[7,:])
print('Predicted y:', y_pred[7])
print('True y: ', y_test.iloc[7])

shap.force_plot(explainer.expected_value,shap_values[7,:],X_test_transformed.
      ↪iloc[7],show = False, matplotlib = True, figsize=(20,3))
plt.savefig(direct + '/figures/Feature_Contribution_Customer7.
      ↪png',dpi=300,bbox_inches = 'tight')
```

tenure	-1.181415
MonthlyCharges	0.490783
TotalCharges	-0.912418


gender_Female	1.000000
gender_Male	0.000000
SeniorCitizen_0	0.000000
SeniorCitizen_1	1.000000
Partner_No	1.000000
Partner_Yes	0.000000
Dependents_No	1.000000
Dependents_Yes	0.000000
PaperlessBilling_No	0.000000
PaperlessBilling_Yes	1.000000
PaymentMethod_Bank transfer (automatic)	0.000000
PaymentMethod_Credit card (automatic)	0.000000
PaymentMethod_Electronic check	1.000000
PaymentMethod_Mailed check	0.000000
PhoneService	1.000000
MultipleLines	2.000000
InternetService	2.000000
OnlineSecurity	1.000000
OnlineBackup	2.000000
DeviceProtection	1.000000
TechSupport	1.000000
StreamingTV	1.000000
StreamingMovies	1.000000
Contract	0.000000

Name: 7, dtype: float64

```
[ 1.88712964 -0.27970072 -0.75400132 -0.09430171  0.06903318  0.1907135
 -0.08514097 -0.06409543  0.09461988 -0.02098952  0.07800132  0.13539609
  0.00208842  0.0188313   0.03479945  0.08634609  0.05023684 -0.06622797
  0.26751503  0.98311442  0.02707661 -0.04589792  0.00775927  0.02428314
 -0.07618488 -0.07185158  0.52258236]
```

Predicted y: 1

True y: 1



```
[916]: # Looking at Costumer No.89
```

```
print(X_test_transformed.iloc[89])
print(shap_values[89,:])
```

```

print('Predicted y:', y_pred[89])
print('True y: ', y_test.iloc[89])

shap.force_plot(explainer.expected_value, shap_values[89], X_test_transformed.
    ↳iloc[89], show=False, matplotlib=True, figsize=(20,3))
plt.savefig(direct + '/figures/Feature_Contribution_Customer89.png', dpi=300,
    ↳bbox_inches = 'tight')

```

tenure	-1.222207
MonthlyCharges	1.034613
TotalCharges	-0.929265
gender_Female	0.000000
gender_Male	1.000000
SeniorCitizen_0	1.000000
SeniorCitizen_1	0.000000
Partner_No	1.000000
Partner_Yes	0.000000
Dependents_No	1.000000
Dependents_Yes	0.000000
PaperlessBilling_No	1.000000
PaperlessBilling_Yes	0.000000
PaymentMethod_Bank transfer (automatic)	0.000000
PaymentMethod_Credit card (automatic)	0.000000
PaymentMethod_Electronic check	1.000000
PaymentMethod_Mailed check	0.000000
PhoneService	1.000000
MultipleLines	1.000000
InternetService	2.000000
OnlineSecurity	1.000000
OnlineBackup	2.000000
DeviceProtection	1.000000
TechSupport	1.000000
StreamingTV	2.000000
StreamingMovies	2.000000
Contract	0.000000


Name: 89, dtype: float64

```

[ 1.94805589 -0.63072667 -0.76657849  0.09620679 -0.07042779 -0.0374127
  0.0167023  -0.06409543  0.09461988 -0.02098952  0.07800132 -0.19974274
 -0.00308094  0.0188313   0.03479945  0.08634609  0.05023684 -0.06622797
 -0.16050902  0.98311442  0.02707661 -0.04589792  0.00775927  0.02428314
  0.30284439  0.31033769  0.52258236]

```

Predicted y: 1
True y: 1



[917]: *# Looking at Costumer No.33*

```
print(X_test_transformed.iloc[33])
print(shap_values[33,:])
print('Predicted y:', y_pred[33])
print('True y: ', y_test.iloc[33])

fig = plt.gcf()
shap.force_plot(explainer.expected_value,shap_values[33],X_test_transformed.
    ↳iloc[33], show=False, matplotlib=True,figsize=(20,3))
plt.savefig(direct + '/figures/Feature_Contribution_Customer33.
    ↳png',dpi=300,bbox_inches = 'tight')
```

tenure	-1.262998
MonthlyCharges	-0.632017
TotalCharges	-0.983465
gender_Female	0.000000
gender_Male	1.000000
SeniorCitizen_0	1.000000
SeniorCitizen_1	0.000000
Partner_No	0.000000
Partner_Yes	1.000000
Dependents_No	0.000000
Dependents_Yes	1.000000
PaperlessBilling_No	0.000000
PaperlessBilling_Yes	1.000000
PaymentMethod_Bank transfer (automatic)	0.000000
PaymentMethod_Credit card (automatic)	0.000000
PaymentMethod_Electronic check	0.000000
PaymentMethod_Mailed check	1.000000
PhoneService	1.000000
MultipleLines	1.000000
InternetService	1.000000
OnlineSecurity	1.000000
OnlineBackup	1.000000

```

DeviceProtection          1.000000
TechSupport              1.000000
StreamingTV              1.000000
StreamingMovies          1.000000
Contract                 0.000000

```

```
Name: 33, dtype: float64
```

```


[ 2.00898214  0.44503286 -0.80704198  0.09620679 -0.07042779 -0.0374127
  0.0167023   0.06915911 -0.10209505  0.04897555 -0.18200307  0.13539609
  0.00208842  0.0188313   0.03479945 -0.04252867 -0.15649914 -0.06622797
 -0.16050902 -0.33827593  0.02707661  0.00861267  0.00775927  0.02428314
 -0.07618488 -0.07185158  0.52258236]

```

```
Predicted y: 0
```

```
True y: 1
```

```
<Figure size 432x288 with 0 Axes>
```



```
[ ]:
```